

---

# **Operating Systems**

## **Spring 2018**

# Syllabus

---

- **Instructors:**

- Dongkun Shin
- Office : Room 85470
- E-mail : dongkun@skku.edu
- Office Hours: Tue. 15:00-17:30 or by appointment

- **Lecture notes**

- [nyx.skku.ac.kr](http://nyx.skku.ac.kr) → Courses → Operation Systems (2018 Spring)
- [http://nyx.skku.ac.kr/?page\\_id=1610](http://nyx.skku.ac.kr/?page_id=1610)
- Lecture notes and talks will be given in **English**.

# Syllabus (cont'd)


---

- **Main text**

- Operating Systems: Three Easy Pieces
- <http://pages.cs.wisc.edu/~remzi/OSTEP/>
- **free online book!!**
- The book is very easy and funny.
- You have to read the scheduled chapter before the class time.

- **Grading policy (subject to change)**

- Attendance: 5% (I'll check at random.)
- Midterm exam: 35%
- Final exam: 35%
- Assignment (Report, Project): 25%
- If you miss one or both of exams, you will fail this course.
- If you aren't in the class room when I call your name, you are regarded as missing the class.
- **Cheating on tests and other assignments will not be tolerated and you will take no (or a negative) point for the test!**



# 출석인정 사유

연번	출석인정 사유	제출서류
1	가족이 사망하여 상중(喪中)인 경우 <배우자, 자녀, 본인 및 배우자의 부모, 조부모(7일 이내), 본인 및 배우자의 형제자매, 백숙부모, 형제자매의 배우자(3일 이내)>	사망진단서 혹은 병원의 확인서
2	학교 공식행사(교내 각 기관이 인정하는 외부기관 행사를 포함한다). 교육실습, 현장수업, 예비군훈련 등 행사(훈련)참석확인서 등으로 결석 사유가 확인되는 경우	행사참석확인서/ 훈련참석확인서 등
3	스포츠단 선수로 스포츠단이 공식인정하는 대회에 참가하는 경우	대회참석확인서
4	학기 중 취업이 확정되어 출석하지 못하는 경우, 해당 수업을 담당하는 교수가 학생에게 그에 상응하는 별도의 추가 과제를 부과하고 이를 성 실히 이행했다고 판단하는 경우	취업확인서 또는 재직증명서 및 출근의무를 입증할 자료 등
5	기타 학장이 부득이 하다고 인정하는 사유가 있는 경우	출석인정사유 확인서

# Syllabus (cont'd)

---

- **Course Outline**

- Part 1. CPU Virtualization
  - Process
  - CPU Scheduling
- Part 2. Memory Virtualization
  - Address Space, Allocation
  - Address Translation
  - Paging, TLB, Swapping
- Part 3. Concurrency
  - Thread
  - Lock, Semaphore
- Part 4. Persistence
  - I/O Systems
  - Storage
  - File System

# Syllabus (cont'd)

---

- **Assignments**

- Homework
  - Reports on advanced topics
  - Homework in the textbook
- Three term projects

- **Prerequisites**

- C programming
- System Programming
- Computer Architecture

---

If you have any questions,  
please feel free to interrupt me  
in English  
or Korean.

---

# **Chap.2**

## **Introduction to Operating Systems**



# What happens when a program runs?

---

- **Von Neumann** model of Computing (Instruction Execution)
  - The processor **fetches** an instruction from memory,
  - **Decodes** it (i.e., figures out which instruction this is), and
  - **Executes** it (i.e., it does the thing that it is supposed to do, like add two numbers together, access memory, check a condition, jump to a function, and so forth).
  - After it is done with this instruction, the processor moves on to the **next** instruction, and so on, and so on

# Operating System (OS)

---

- Makes it easy to run programs
  - even many programs at the same time
- Allows programs to share memory
- Enables programs to interact with devices
- ➔ in charge of making sure the system operates correctly and efficiently in an easy-to-use manner

# Operating System (OS)

---

- The role of OS

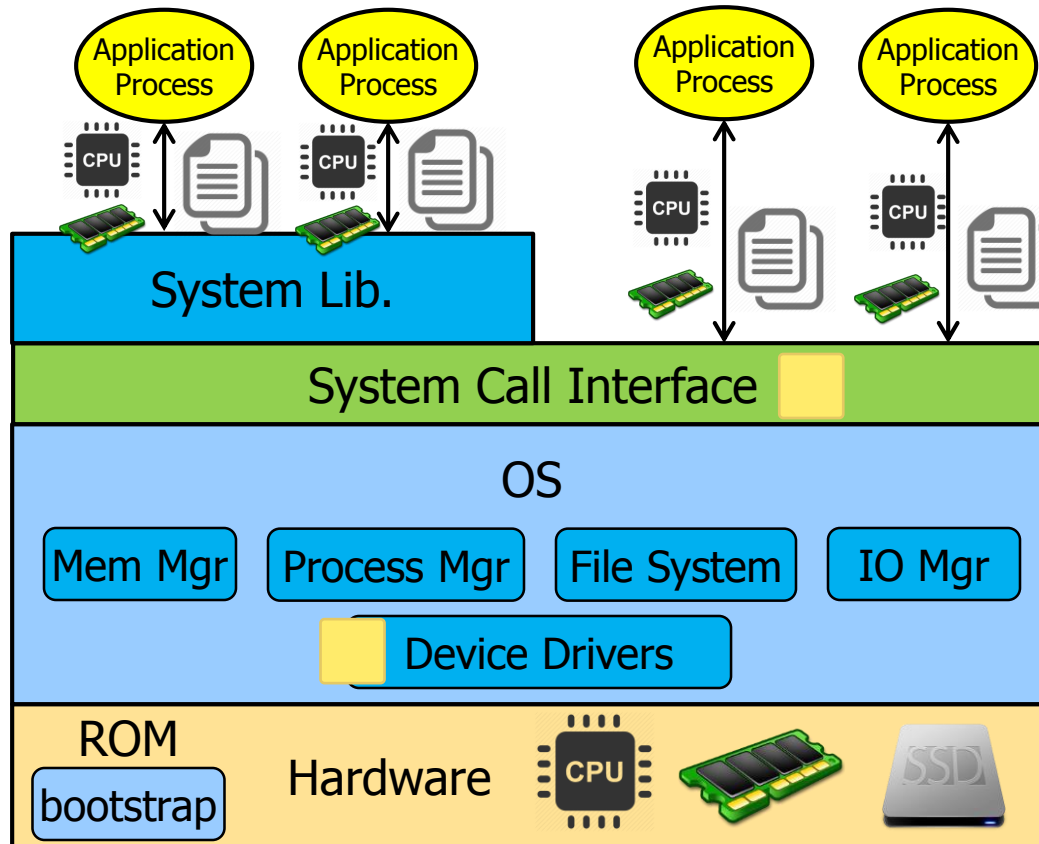
- Virtualization

- OS takes a physical resource (such as the processor, or memory, or a disk)
    - transforms it into a more general, powerful, and easy-to-use virtual form of itself.
    - System calls allow users to tell the OS what to do and thus make use of the features of the OS
    - Virtualization allows many programs to run concurrently

- Resource Manager

- Resource: CPU, memory, and disk
    - OS manages those resources efficiently or fairly or indeed with many other possible goals in mind.

# Computer System Organization



# Virtualizing the CPU

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

int main(int argc, char *argv[])
{
    if (argc != 2) {
        fprintf(stderr, "usage: cpu <string>\n");
        exit(1);
    }
    char *str = argv[1];

    while (1) {
        printf("%s\n", str);
        Spin(1);
    }
    return 0;
}
```

Spin(1): repeatedly checks the time and returns once it has run for a second

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
A
^C
prompt>
```

```
prompt> ./cpu A & ; ./cpu B & ; ./cpu C & ; ./cpu D &
[1] 7353
[2] 7354
[3] 7355
[4] 7356
```

A  
B  
D  
C  
A  
B  
D  
C  
A  
C  
B  
D  
...

*Even though we have  
only one processor,  
somehow all four of these  
programs seem to be  
running at the same time!*

Running Many Programs At Once

# Virtualizing the CPU

---

- Illusion that the system has a very **large number of virtual CPUs**.
- Turning a single CPU (or small set of them) into a seemingly infinite number of CPUs and thus allowing many programs to seemingly run at once
- **Two Questions**
  - Policy
    - If more than one program want to run at a particular time, which *should* run?
  - Mechanisms
    - How to implement the ability to run multiple programs at once?

# Virtualizing Memory

```
int
main(int argc, char *argv[])
{
    int *p;
    p = malloc(sizeof(int));
    assert(p != NULL);
    printf("(%)d) address pointed to by p: %pn", getpid(), p);
    *p = 0;
    while (1) {
        Spin(1);
        *p = *p + 1;
        printf("(%)d) p: %d\n", getpid(), *p);
    }
    return 0;
}
```

*Each running program has allocated memory at the same address (0x200000), and yet each seems to be updating the value at 0x200000 independently!*

*It is as if each running program has its **own private memory**, instead of sharing the same physical memory with other running programs*


```
prompt> ./mem
(2134) address pointed to by p: 0x200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

```
prompt> ./mem &; ./mem &
[1] 24113
[2] 24114
(24113) address pointed to by p: 0x200000
(24114) address pointed to by p: 0x200000
(24113) p: 1
(24114) p: 1
(24114) p: 2
(24113) p: 2
(24113) p: 3
(24114) p: 3
(24113) p: 4
(24114) p: 4
...
```

# Virtualizing Memory

---

- Each process accesses its own private **virtual address space**
- OS maps the virtual address onto the physical memory of the machine.
- A memory reference within one running program does not affect the address space of other processes (or the OS itself); as far as the running program is concerned, it has physical memory all to itself.
- The reality, however, is that physical memory is a shared resource, managed by the operating system.





# Concurrency

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"

volatile int counter = 0;
int loops;

void *worker(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        counter = counter + 1;
    }
    pthread_exit(NULL);
}

int main(int argc, char *argv[]) {
    if (argc != 2) {
        fprintf(stderr, "usage: threads <loops>\n");
        exit(1);
    }
    loops = atoi(argv[1]);
    pthread_t p1, p2;
    printf("Initial value : %d\n", counter);
    Pthread_create(&p1, NULL, worker, NULL);
    Pthread_create(&p2, NULL, worker, NULL);
    Pthread_join(p1, NULL);
    Pthread_join(p2, NULL);
    printf("Final value : %d\n", counter);
    return 0;
}
```

*multi-threaded programs*

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

**counter = counter + 1; → no atomic**  
**3 instructions**

- (1) load the value of the counter from memory into a register
- (2) increment it
- (3) store it back into memory.

**How can we build a correctly working multi-threaded program?**  
**What primitives are needed from the OS?**

# Persistence

- DRAM is **volatile**
- We need hardware and software to store data **persistently**
  - H/W: HDD, SSD
  - S/W: **file system** manages the disk, responsible for storing any files the user creates in a **reliable** and **efficient** manner on the disks of the system

```
#include <stdio.h>
#include <unistd.h>
#include <assert.h>
#include <fcntl.h>
#include <sys/types.h>


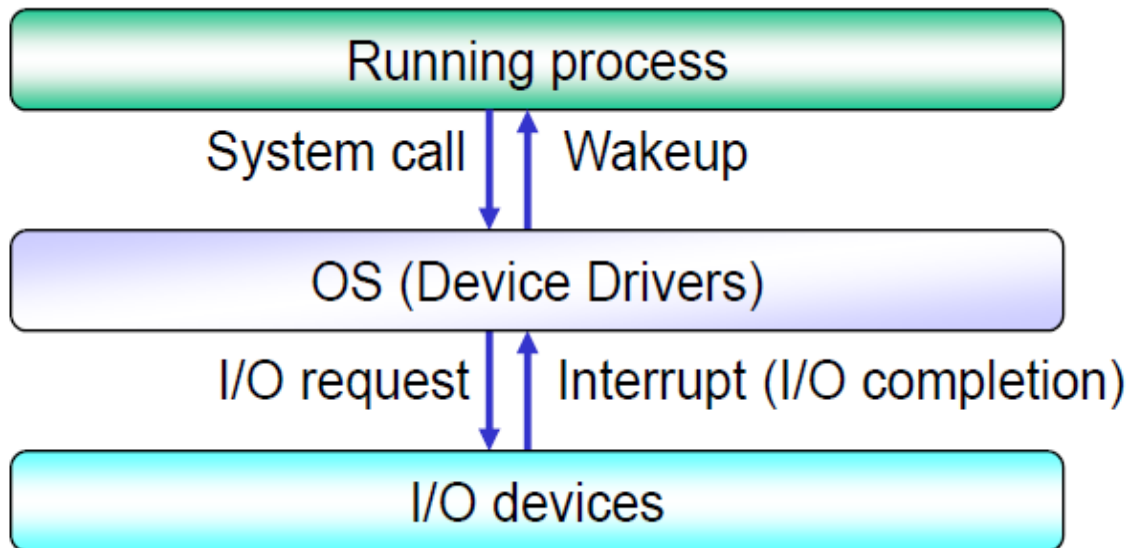
int
main(int argc, char *argv[])
{
    int fd = open("/tmp/file", O_WRONLY | O_CREAT |
                  O_TRUNC, S_IRWXU);
    assert(fd > -1);
    int rc = write(fd, "hello world\n", 13);
    assert(rc == 13);
    close(fd);
    return 0;
}
```

- OS does not create a private, virtualized disk for each application.
- Rather, users will want to share information that is in files.
- System calls: open, read, write, close
- Device driver issues I/O requests to the underlying storage device

# Persistence

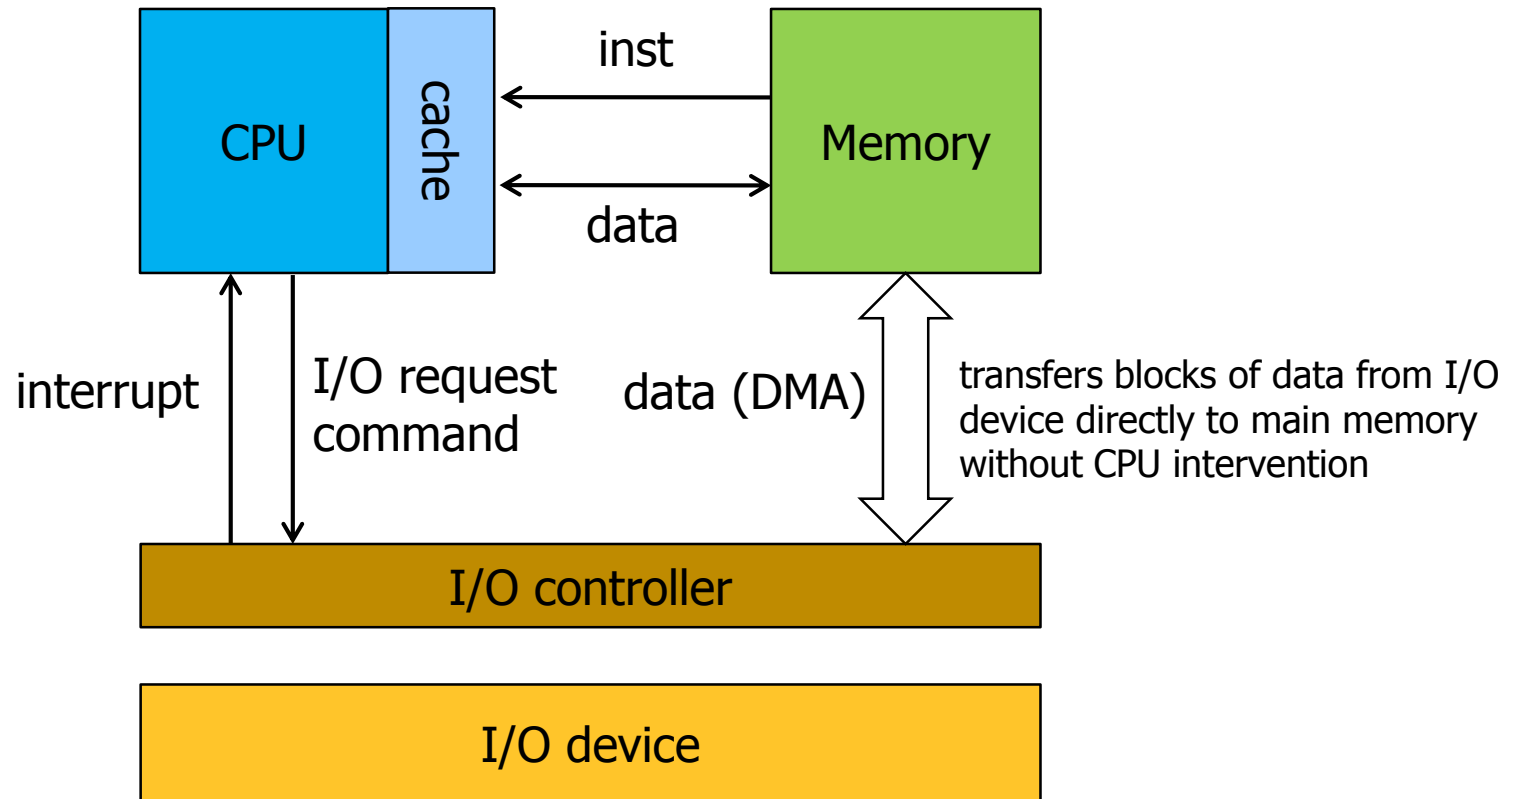
---

- **I/O management**



# I/O Mechanism

---



OS includes interrupt handlers

# Design Goals

---

- build up some **abstractions** in order to make the system convenient and easy to use
- provide high **performance**
- minimize the **overheads** of the OS
  - Extra time and space
- provide **protection** between applications, as well as between the OS and applications
  - isolating processes from one another is the key to protection
- **Reliability**
  - The operating system must also run non-stop
  - when it fails, all applications running on the system fail as well
- **energy-efficiency, security, mobility**
- Depending on how the system is used, the OS will have different goals and thus likely be implemented in at least slightly different ways.

# History

---

- **Early Operating Systems: Just Libraries**
- **Beyond Libraries: Protection**
  - The idea of a **system call** was invented
  - System call transfers control (i.e., jumps) into the OS while simultaneously raising the **hardware privilege level**.
  - User applications run in what is referred to as **user mode** which means the hardware restricts what applications can do
  - **Trap** raises the privilege level to **kernel mode**, the OS has full access to the hardware of the system
- **The Era of Multiprogramming, Minicomputer**
  - OS loads a number of jobs into memory and switch rapidly between them, thus improving CPU utilization
  - Protection mechanisms are necessary to control access to system resources (including files)
  - **concurrency** issues
  - The introduction of the **UNIX** operating system
    - Ken Thompson (and Dennis Ritchie) at Bell Labs

# History

---

- **The Modern Era**

- personal computer: Apple II, IBM PC
  - Unfortunately, for operating systems, the PC at first represented a great leap backwards
  - forgot (or never knew of) the lessons learned in the era of minicomputers
  - DOS: no memory protection
  - Mac OS (v9 and earlier): cooperative job scheduling
- Now
  - Mac OS X
    - Steve Jobs took his UNIX-based NeXTStep operating environment with him to Apple
  - Windows NT
  - Linux
    - **Linus Torvalds** wrote his own version of UNIX which borrowed heavily on the principles and ideas behind the original system, but not from the code base, thus avoiding issues of legality

# Homework

---

- **Submit a report on “The Evolution of the Unix Time-sharing System”**