

---

# **Page Replacement**

## **Chap 21, 22**

# Virtual Memory Concept

---

- **Virtual memory**

- Concept

- A technique that allows the execution of processes that are not completely in memory

- Partition each user's program into multiple blocks

- Load into memory the blocks that is necessary at each time during execution

- Only part of the program needs to be in memory for execution

- Noncontiguous allocation

- Logical memory size is not constrained by the amount of physical memory that is available

- Separation of logical memory as perceived by users from physical memory

# Virtual Memory Concept

---

- **Virtual memory**

- Benefits

- Easier programming

- Programmer no longer needs to worry about the amount of physical memory available

- Allows address spaces to be shared by several processes

- Higher multiprogramming degree

- Increase in CPU utilization and throughput (not in response time and turnaround time)

- Less I/O for loading and swapping processes into memory

- Faster execution of processes


# Virtual Memory Concept

---

- **Virtual memory**
  - Drawbacks
    - Address mapping overhead
    - Page fault handling overhead
    - Not adequate for real-time (embedded) systems

# Virtual Memory That is Larger Than Physical Memory

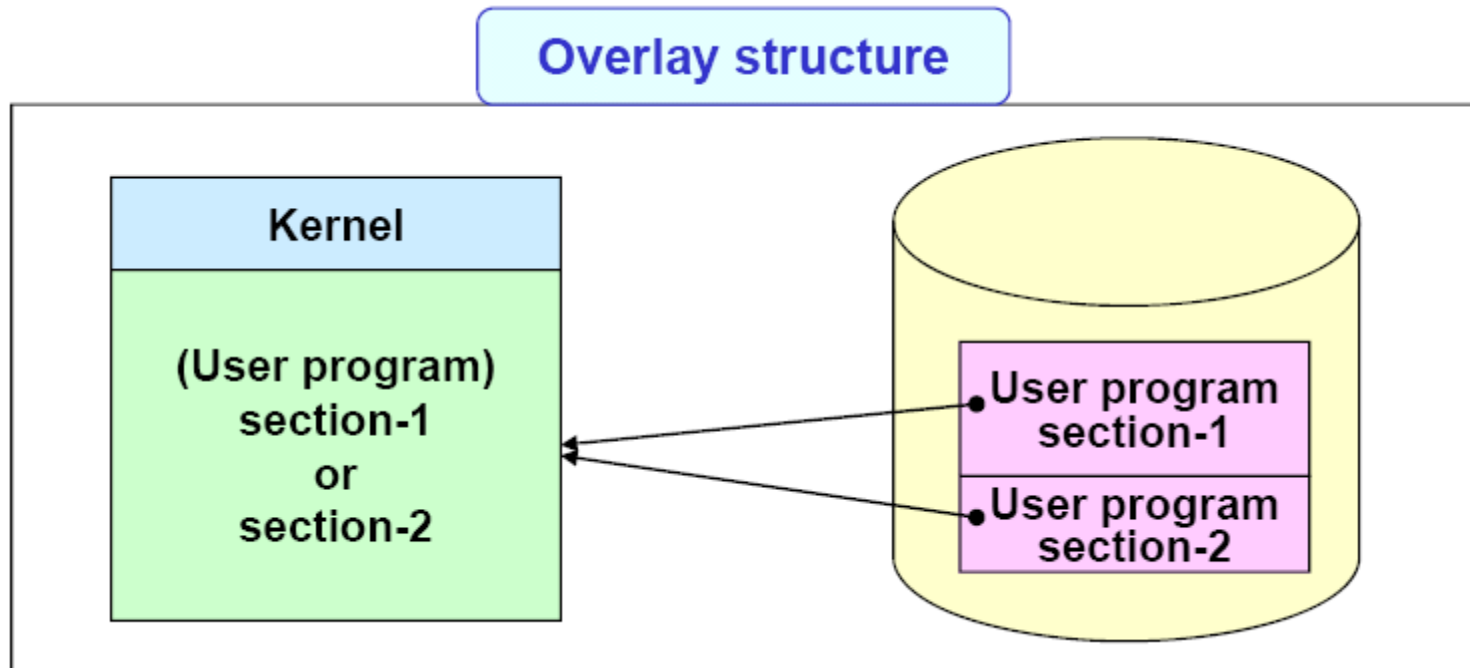
**OS make use of a larger, slower device to transparently provide the illusion of a large virtual address space**



# Memory Overlay (old system)

---


- **Program-size > memory-size**
  - w/o OS support
  - Requires support from compiler/linker/loader



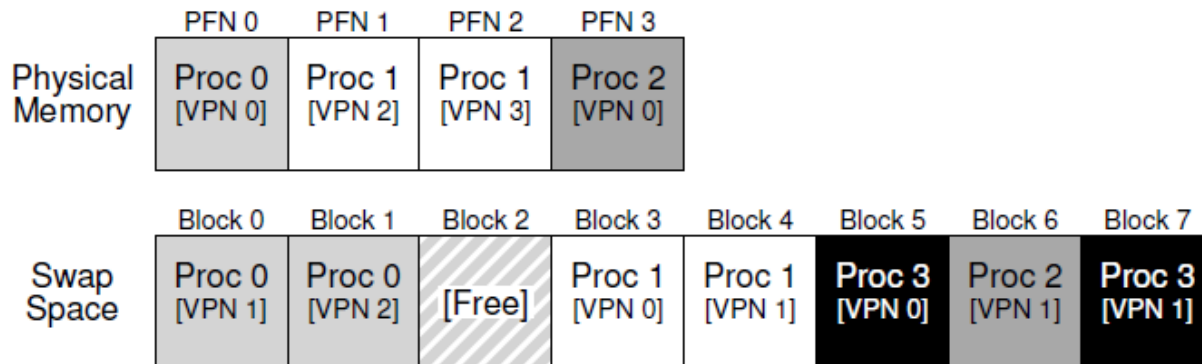
# Swapping

- A process can be swapped temporarily out of memory to a **backing store (swap device)**

Process-level  
swapping



Page-level  
swapping



# Swapping

---

- **Notes on swapping**

- Time quantum vs swap time
  - Time quantum should be substantially larger than swap time (context switch time) for efficient CPU utilization
- Memory areas to be swapped out
  - Swap only what is actually used
- Pending I/O
  - If the I/O is asynchronously accessing the user memory for I/O buffers, then the process cannot be swapped
  - Solutions
    - **Never swap a process with pending I/O**
    - Execute I/O operations only into kernel buffers (and deliver it to the process memory when the process is swapped in)



# Demand Paging

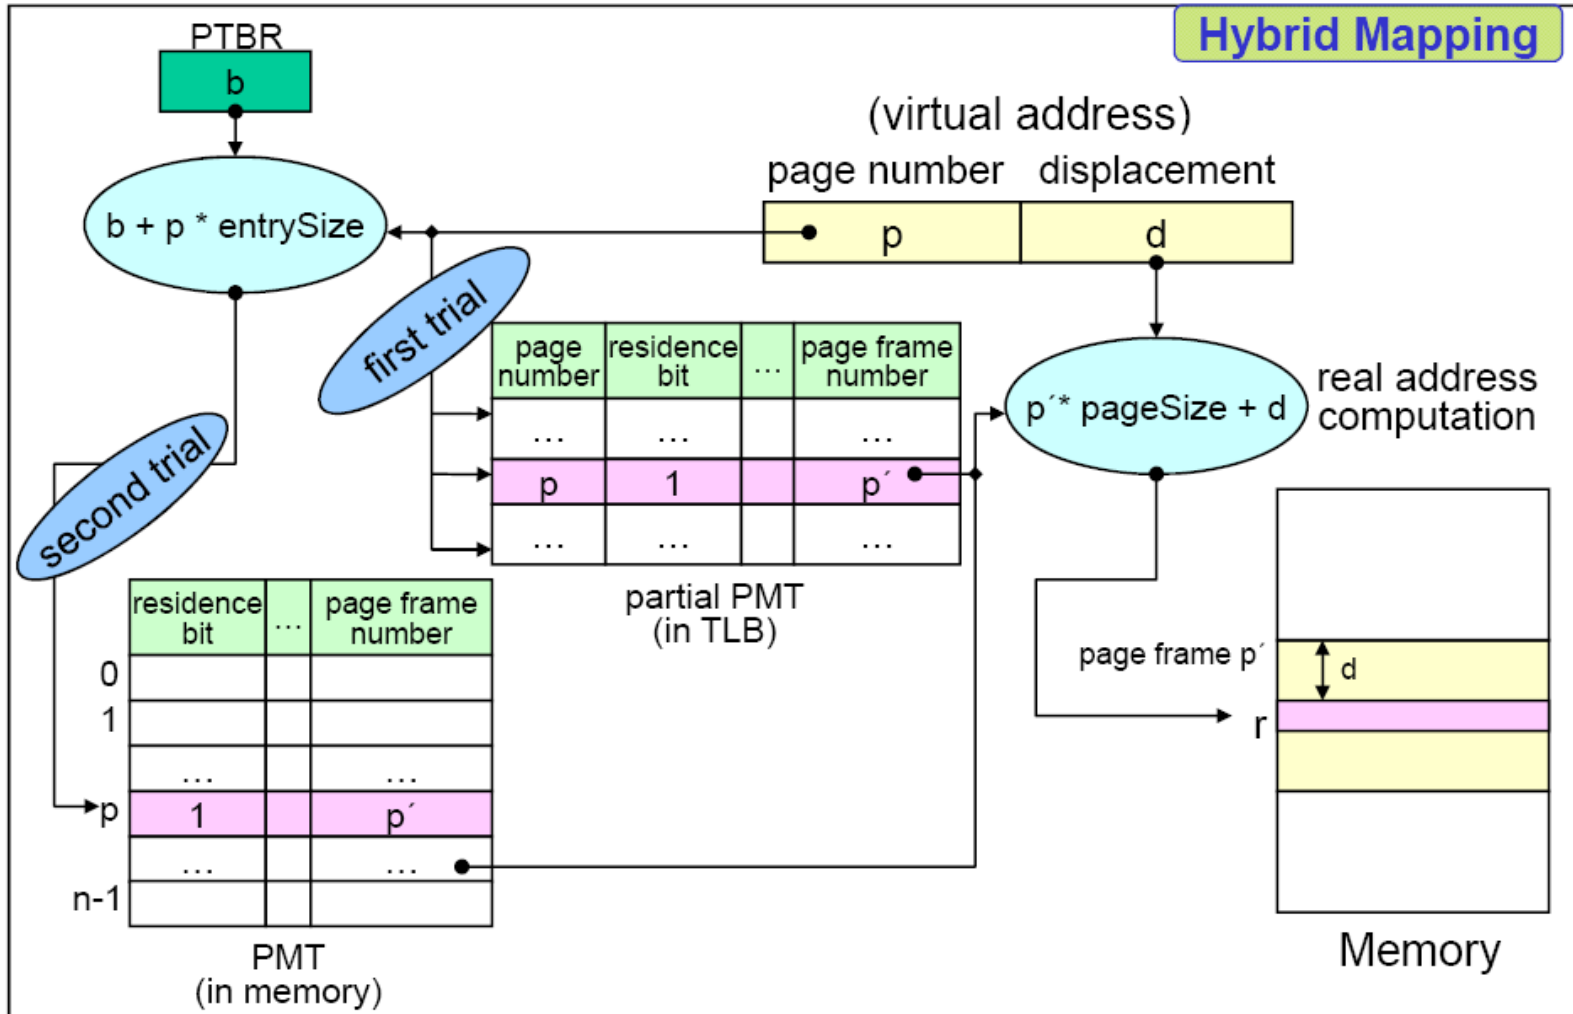
---

- **Paging (Demand paging) system**

- Partition the program into the same size blocks (pages)
- Loading of executable program
  - Initially, load pages only as they are needed
  - During execution, load the pages when they are demanded (referenced)
  - Pages that are never accessed are never loaded into physical memory
- With each page table entry a **present (residence)** bit is associated
  - **Present = true:** in-memory, **memory resident**
  - **Present = false:** not-in-memory
- Initially present bit is set to **false** on all entries
- During MMU address translation, if present bit in page table entry is **false**  $\Rightarrow$  page fault

# Demand Paging

- Address Mapping



# Page Fault

---

- If there is a reference to a page, first reference to that page will trap to operating system:  
**page fault**
  1. Operating system looks at another table to decide:
    - Invalid reference  $\Rightarrow$  abort
    - Just not in memory
  2. Find free frame
  3. Swap page into frame via scheduled disk operation
  4. Reset tables to indicate page now in memory  
Set **present** bit = **T**
  5. Restart the instruction that caused the page fault

# Steps in Handling a Page Fault

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True)    // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset    = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10 else                    // TLB Miss
11     PTEAddr = PTBR + (VPN * sizeof(PTE))
12     PTE = AccessMemory(PTEAddr)
13     if (PTE.Valid == False)
14         RaiseException(SEGMENTATION_FAULT)
15     else
16         if (CanAccess(PTE.ProtectBits) == False)
17             RaiseException(PROTECTION_FAULT)
18         else if (PTE.Present == True)
19             // assuming hardware-managed TLB
20             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
21             RetryInstruction()
22         else if (PTE.Present == False)
23             RaiseException(PAGE_FAULT)
```

```
1  PFN = FindFreePhysicalPage()
2  if (PFN == -1)                // no free page found
3      PFN = EvictPage()         // run replacement algorithm
4  DiskRead(PTE.DiskAddr, pfn) // sleep (waiting for I/O)
5  PTE.present = True           // update page table with present
6  PTE.PFN      = PFN           // bit and translation (PFN)
7  RetryInstruction()           // retry instruction
```

# Stages in Demand Paging

---

- 1. Trap to the operating system**
- 2. Save the user registers and process state**
- 3. Determine that the interrupt was a page fault**
- 4. Check that the page reference was legal and determine the location of the page on the disk**
- 5. Issue a read from the disk to a free frame:**
  1. Wait in a queue for this device until the read request is serviced
  2. Wait for the device seek and/or latency time
  3. Begin the transfer of the page to a free frame
- 6. While waiting, allocate the CPU to some other user**
- 7. Receive an interrupt from the disk I/O subsystem (I/O completed)**
- 8. Save the registers and process state for the other user**
- 9. Determine that the interrupt was from the disk**
- 10. Correct the page table and other tables to show page is now in memory**
- 11. Wait for the CPU to be allocated to this process again**
- 12. Restore the user registers, process state, and new page table, and then resume the interrupted instruction**

# Performance of Demand Paging

---

- **Effective access time**

- Memory access time
  - 10 ~ 200 nanoseconds (Assume 200ns)
- Average paging service time: about 8 ms
- Page fault rate:  $p$  ( $0 \leq p \leq 1$ )
- EAT(Effective Access Time)
  - $$\begin{aligned} \text{EAT} &= (1-p) \cdot m_a + p \cdot \text{PagingTime} \\ &= (1-p) \cdot 200 + p \cdot 8,000,000 \\ &= 200 + 7,999,800 \cdot p \end{aligned}$$
  - When  $p = 1/1000$ ,  $\text{EAT} = 8.2 \text{ us}$  ( $40 \times m_a$ )
- If we want less than 10% degradation,
  - $\text{EAT} = 200 + 7,999,800 \cdot p < 220$
  - $P < 0.0000025$  ( $= 1/400,000$ )

# Demand Paging Optimizations

---

- **Swap space I/O faster than file system I/O even if on the same device**
  - Swap allocated in larger chunks, less management needed than file system
- **Demand page in from program binary on disk, but discard rather than paging out when freeing frame**
  - Still need to write to swap space
    - Pages not associated with a file (like stack and heap) – **anonymous memory**
    - Pages modified in memory but not yet written back to the file system
- **Prefetching**
  - OS could guess that a page is about to be used, and thus bring it in ahead of time
- **Mobile systems**
  - Typically don't support swapping
  - Instead, demand page from file system and reclaim read-only pages (such as code)
  - cf. **zswap**

# Page Replacement

---

- Prevent over-allocation of memory
- Use **modify** (**update**, **dirty**) bit to reduce overhead of page transfers
  - only modified pages are written to disk
  - If `modify == 1`, the contents of the page in memory and in disk are not same
    - Write-back (to disk) is necessary for the page
- Large virtual memory can be provided on a smaller physical memory



# Page Replacement


---

1. Find the location of the desired page on disk
2. Find a free frame:
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a **victim frame**
    - Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap

Note now potentially 2 page transfers for page fault – increasing EAT

# Graph of Page Faults Versus The Number of Frames

---



# When Replacements Really Occur?

---

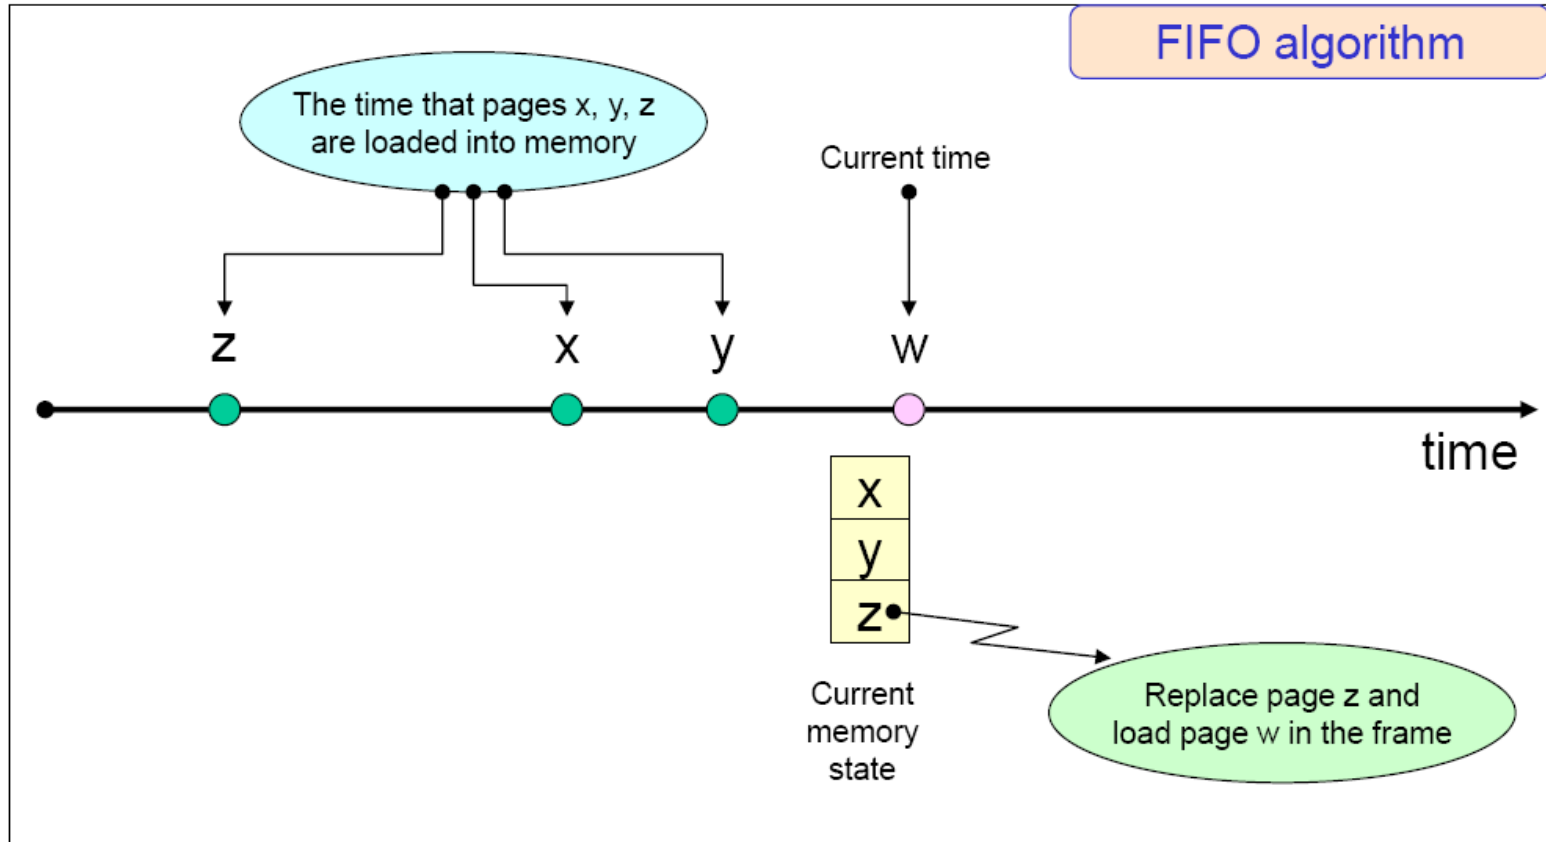
- OS keeps a small portion of memory free more proactively
- **Watermark** scheme
  - high watermark (HW) and low watermark (LW)
  - When OS notices that there are fewer than LW pages available, a background thread (**swap daemon** or **page daemon**) that is responsible for freeing memory runs.
  - The thread evicts pages until there are HW pages available.
  - The background thread then goes to sleep.
  - many systems will **cluster** or **group** a number of pages and write them out at once to the swap partition, thus increasing the efficiency of the disk

# First-In-First-Out (FIFO) Algorithm

---

- Choose the page to be replaced based on when the page is previously loaded into memory
- Scheme
  - Replace the oldest page
- Requirements
  - Timestamping (memory load time for each page) is necessary
- Characteristics
  - May replace frequently used pages
- **FIFO anomaly (Belady's anomaly)**
  - In FIFO algorithm, page fault frequency may increase even if more memory frames are allocated

# First-In-First-Out (FIFO) Algorithm



# First-In-First-Out (FIFO) Algorithm

- FIFO algorithm: Example**

- 4 page frames allocated, initially empty

$\omega = 1\ 2\ 6\ 1\ 4\ 5\ 1\ 2\ 1\ 4\ 5\ 6\ 4\ 5$

Memory state change (FIFO)

Time	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Ref. string	1	2	6	1	4	5	1	2	1	4	5	6	4	5
Memory state	1	1	1	1	1	5	5	5	5	5	5	5	4	4
		2	2	2	2	2	1	1	1	1	1	1	1	5
			6	6	6	6	6	2	2	2	2	2	2	2
					4	4	4	4	4	4	4	6	6	6
Page fault	F	F	F		F	F	F	F				F	F	F

▪ Number of page faults: 10

# First-In-First-Out (FIFO) Algorithm

- FIFO algorithm: Anomaly example

$\omega = 1\ 2\ 3\ 4\ 1\ 2\ 5\ 1\ 2\ 3\ 4\ 5$

Number of page frames: 3

Time	1	2	3	4	5	6	7	8	9	10	11	12
Ref. string	1	2	3	4	1	2	5	1	2	3	4	5
Memory state	1	1	1	4	4	4	5	5	5	5	5	5
		2	2	2	1	1	1	1	1	3	3	3
			3	3	3	2	2	2	2	2	4	4
Page fault	F	F	F	F	F	F	F			F	F	

▪ Number of page faults: 9

Number of page frames: 4

Time	1	2	3	4	5	6	7	8	9	10	11	12
Ref. string	1	2	3	4	1	2	5	1	2	3	4	5
Memory state	1	1	1	1	1	1	5	5	5	5	4	4
		2	2	2	2	2	2	1	1	1	1	5
			3	3	3	3	3	3	2	2	2	2
			4	4	4	4	4	4	4	3	3	3
Page fault	F	F	F	F				F	F	F	F	F

▪ Number of page faults: 10

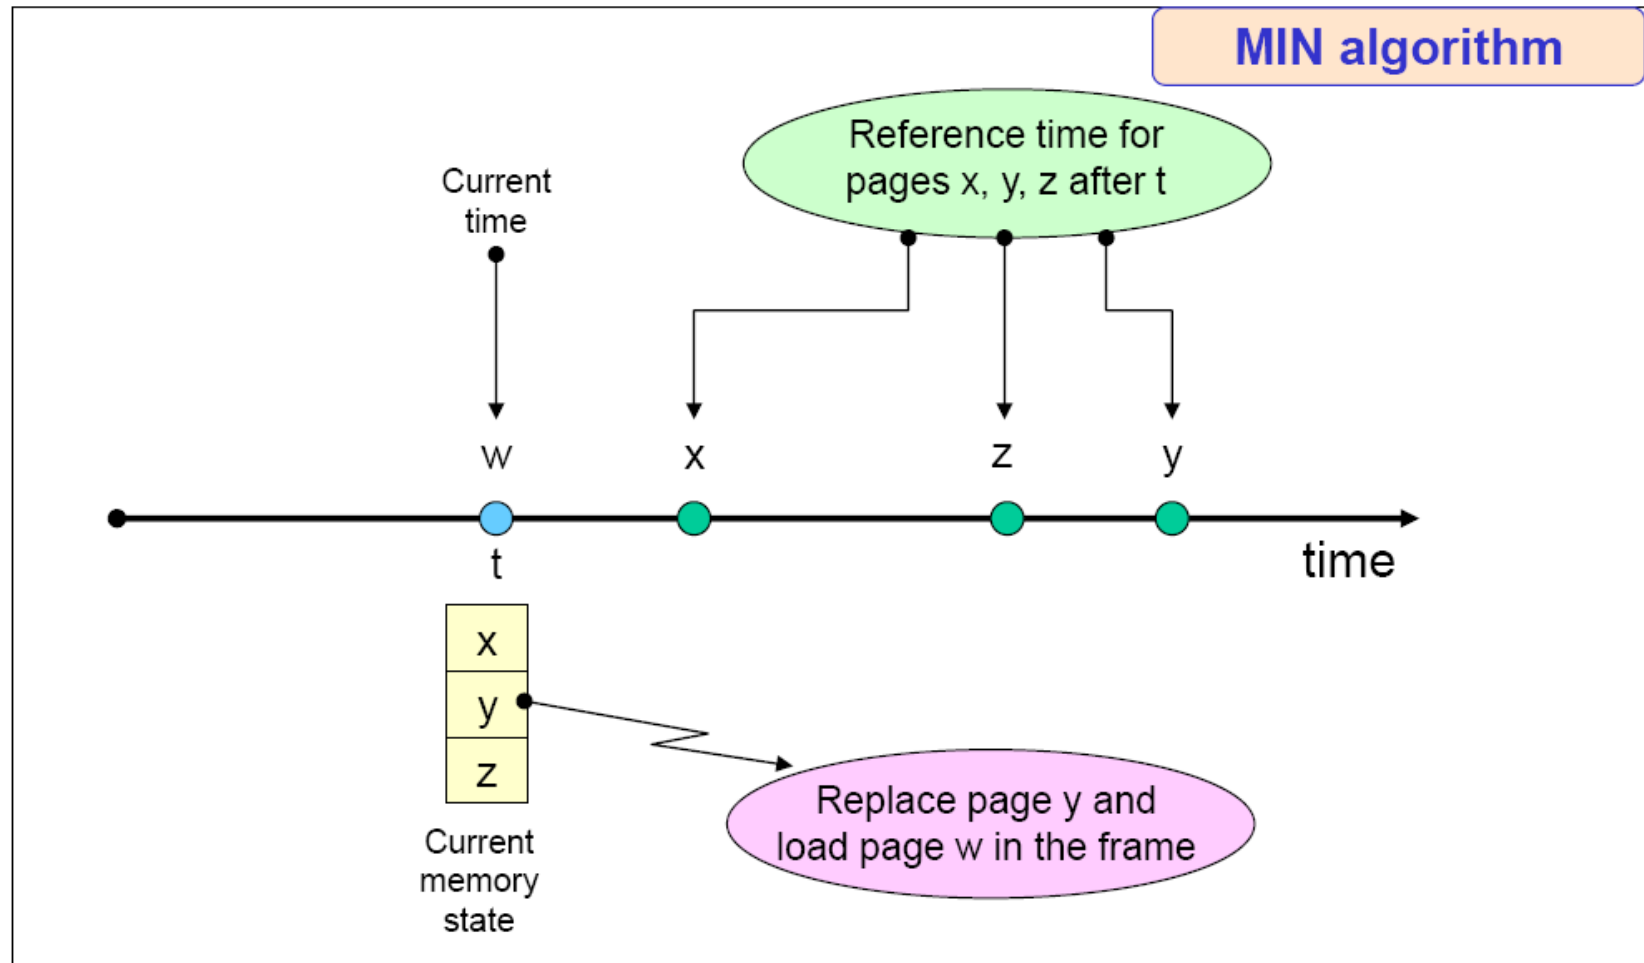
# MIN algorithm (OPT algorithm)

---

- Proposed by Belady in 1966
- Minimizes page fault frequency (proved)
- Scheme
  - Replace the page that will not be used for the longest period of time
  - Tie-breaking rule
    - Page with greatest (or smallest) page number
- Unrealizable
  - Can be used only when the process's reference string is known a priori
- Usage
  - Performance measurement tool for replacement schemes



# MIN algorithm (OPT algorithm)



# MIN algorithm (OPT algorithm)

- **MIN algorithm: Example**

- 4 page frames allocated, initially empty

$\omega = 1\ 2\ 6\ 1\ 4\ 5\ 1\ 2\ 1\ 4\ 5\ 6\ 4\ 5$

Memory state change (MIN)

Time	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Ref. string	1	2	6	1	4	5	1	2	1	4	5	6	4	5
Memory state	1	1	1	1	1	1	1	1	1	1	1	6	6	6
		2	2	2	2	2	2	2	2	2	2	2	2	2
			6	6	6	5	5	5	5	5	5	5	5	5
					4	4	4	4	4	4	4	4	4	4
Page fault	F	F	F		F	F						F		

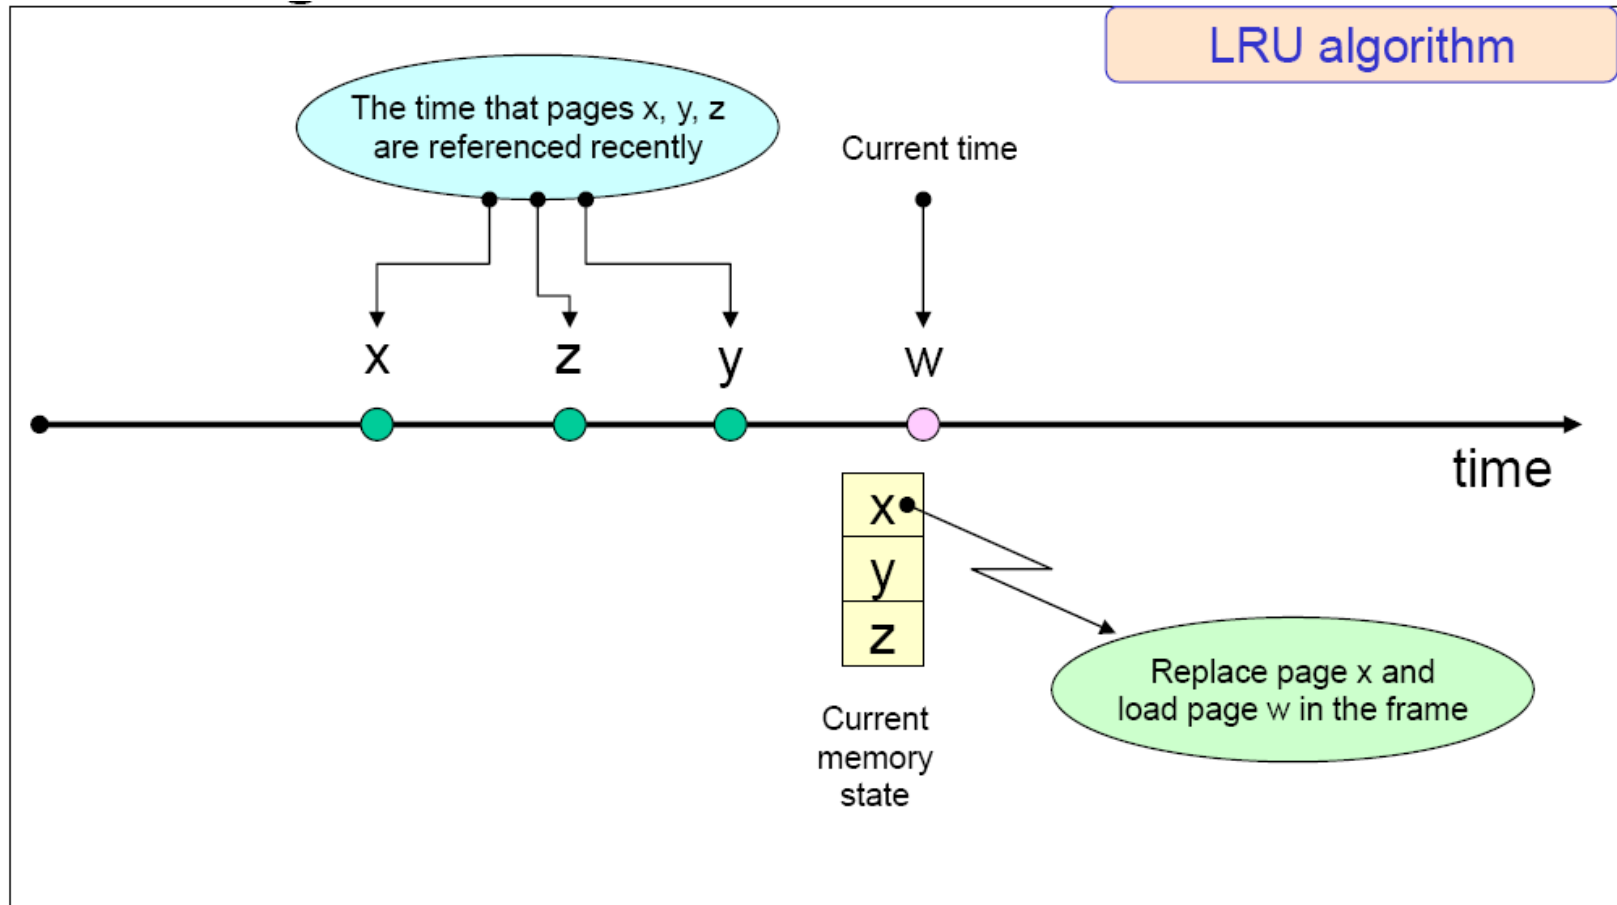
▪ Number of page faults: 6

# Least Recently Used (LRU) Algorithm

---

- Choose the page to be replaced based on the reference time
- Scheme
  - Replace the page that has not been used for the longest period of time
- Requirements
  - Timestamping (page reference time) is necessary
- Characteristics
  - Based on program locality
  - Approximates to the performance of MIN algorithm
- **Used in most practical systems**
- Drawbacks
  - Timestamping overhead at every page reference
  - Number of page faults increases steeply when the process executes large loop with insufficiently allocated memory

# Least Recently Used (LRU) Algorithm



# Least Recently Used (LRU) Algorithm

- LRU algorithm: Example

- 4 page frames allocated, initially empty

$\omega = 1\ 2\ 6\ 1\ 4\ 5\ 1\ 2\ 1\ 4\ 5\ 6\ 4\ 5$

Memory state change (LRU)

Time	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Ref. string	1	2	6	1	4	5	1	2	1	4	5	6	4	5
Memory state	1	1	1	1	1	1	1	1	1	1	1	1	1	1
		2	2	2	2	5	5	5	5	5	5	5	5	5
			6	6	6	6	6	2	2	2	2	6	6	6
					4	4	4	4	4	4	4	4	4	4
Page fault	F	F	F		F	F		F				F		

▪ Number of page faults: 7

# Implementation of LRU algorithm

---

- By counter
  - Use PMT with count field
  - Increment processor clock or counter for each memory access
  - Record the value of processor clock or counter in the corresponding PMT entry for each page reference
  - Can get the relative order of recent access to each page
  - PMT search for selecting a page to be replaced

# Implementation of LRU algorithm

---

- By stack
  - Stack
    - Stack for each process, whose entry is page number
    - Maintains the stack elements (page numbers) in the order of recent access
    - Can delete an element in the middle of the stack
  - When no page fault
    - Deletes the referenced page number from the stack, and inserts it on top of the stack
  - When page fault
    - Displaces the page whose number is at the bottom of the stack, deletes it from the stack, and inserts incoming page number on top of the stack

# LRU Approximation Algorithms

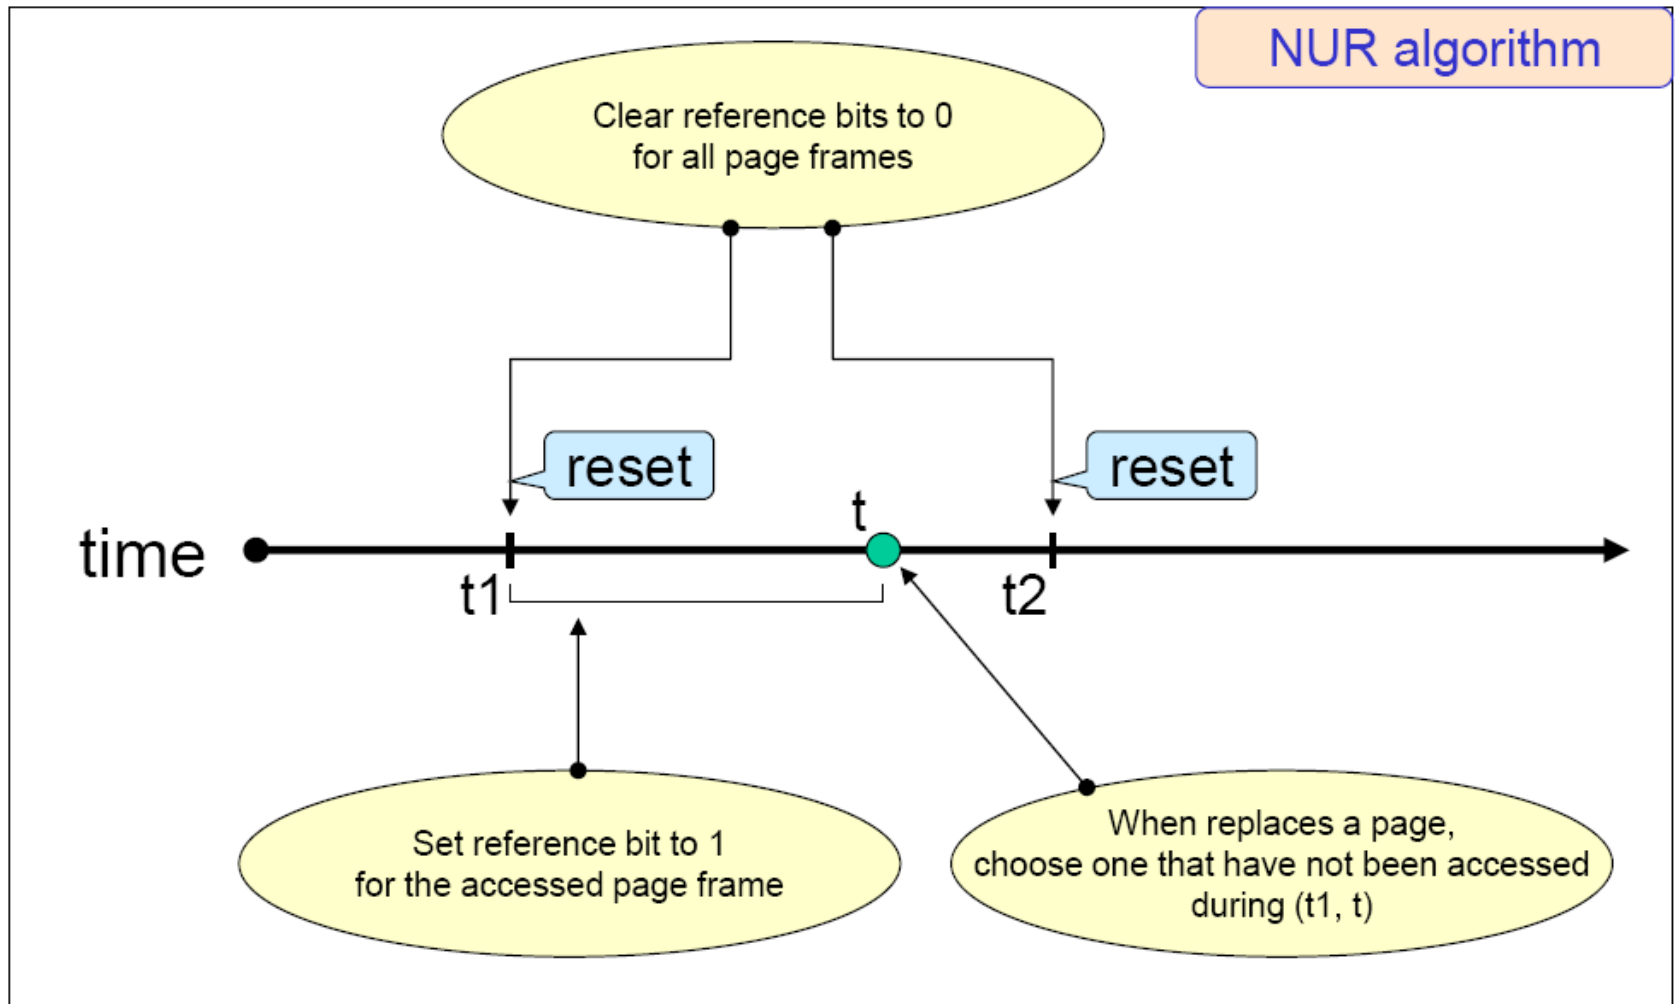
---

- **LRU needs special hardware and still slow**
- **NUR(Not Used Recently) algorithm**
  - LRU approximation scheme
  - Lower overhead than LRU algorithm
  - Uses bit vectors
    - Reference bit vector
    - Update bit vector
  - Scheme
    - Check reference/update bit and choose a victim page
    - Order of replacement (reference bit:  $r$ , update bit:  $m$ )
      - ① Replace the page with  $(r, m) = (0, 0)$
      - ② Replace the page with  $(r, m) = (0, 1)$
      - ③ Replace the page with  $(r, m) = (1, 0)$
      - ④ Replace the page with  $(r, m) = (1, 1)$



# LRU Approximation Algorithms

- **NUR algorithm**



# LRU Approximation Algorithms

---

- **Additional reference-bits algorithm**
  - LRU approximation
  - History register for each page
  - Recording the reference bits at regular intervals
    - Timer interrupts at regular intervals
    - OS shifts the reference bit for each page into the high-order bit of its history register, shifting the other bits right by one bit and discarding the low-order bit
  - Replacement based on the value of history register
    - Interpret the value of the history register as unsigned integers
    - Choose the page with smallest value as a victim

# LRU Approximation Algorithms


---

- **Additional reference-bits algorithm**
- Example
  - Value of the 8-bit history registers
    - 00000000
      - No reference during 8 time periods
    - 11111111
      - Referenced at least once in each period
    - Page a with 01100100 and page b with 00110111
      - Page a has been used more recently than page b

# LRU Approximation Algorithms

---

- **Clock algorithm (Second-chance algorithm)**
  - Used for IBM VM/370 OS
  - Uses reference bit
    - No periodical reset for reference bits
  - Choose the page to be replaced using pointer that circulates the list of pages (page frames) in memory



# LRU Approximation Algorithms

---

- **Clock algorithm**

- Scheme

- Choose the page to be replaced with the pointer moving clockwise
    - Mechanism
      - ① Check the reference bit  $r$  of the page that the pointer points
      - ② If  $r == 0$ , select the page as a victim
      - ③ If  $r == 1$ , reset the reference bit to 0 and advances the pointer goto step-①

- Characteristics

- Earlier memory load time → higher probability of displacement
      - Similar to FIFO algorithm
    - Page replacement based on the reference bit
      - Similar to LRU (or NUR) algorithm

# LRU Approximation Algorithms

## • Clock algorithm: Example

### ▪ Assumptions

- 4 page frames are allocated to the process
- Initially, it has pages a, b, c, d in memory
- Reference bits of the 4 page frames are all 1

$\omega = c a d b e b a b c d$

### Memory state change (Clock)

Time		0	1	2	3	4	5	6	7	8	9	10
Ref. string			c	a	d	b	e	b	a	b	c	d
Memory state	frame 0		→a/1	→a/1	→a/1	→a/1	e/1	e/1	e/1	e/1	→e/1	d/1
	frame 1		b/1	b/1	b/1	b/1	→b/0	→b/1	b/0	b/1	b/1	→b/0
	frame 2		c/1	c/1	c/1	c/1	c/0	c/0	a/1	a/1	a/1	a/0
	frame 3		d/1	d/1	d/1	d/1	d/0	d/0	→d/0	→d/0	c/1	c/0
Page fault							F		F		F	F
Pclock							e		a		c	d
Qclock							a		c		d	e

- Pclock : Pages loaded into memory
- Qclock : Pages displaced from memory

# LRU Approximation Algorithms

---

- **Enhanced clock algorithm**

- Similar to clock algorithm
- Considers **update (dirty, modified) bit** as well as reference bit

- Scheme

- Choose the page to be replaced with the pointer moving clockwise
- Mechanism

- ① Check  $(r, m)$  of the page that the pointer points
- ②  $(0, 0)$ : select the page as a victim, advances the pointer
- ③  $(0, 1)$ : set  $(r, m)$  to  $(0, 0)$ , put the page on the cleaning list, goto step-⑥
- ④  $(1, 0)$ : set  $(r, m)$  to  $(0, 0)$ , goto step-⑥
- ⑤  $(1, 1)$ : set  $(r, m)$  to  $(0, 1)$ , goto step-⑥
- ⑥ Advances the pointer, goto step-①

# LRU Approximation Algorithms

## Enhanced clock algorithm: Example

### Assumptions

- 4 page frames are allocated to the process
- Initially, it has pages a, b, c, d in memory
- All reference bits are 1, all update bits are 0

$\omega = c a^W d b^W e b a^W b c d$

### Memory state change (Enhanced clock)

Time		0	1	2	3	4	5	6	7	8	9	10
Ref. string			c	$a^W$	d	$b^W$	e	b	$a^W$	b	c	d
Memory state	frame 0	→a/10	→a/10	→a/11	→a/11	→a/11	<u>a/00</u>	<u>a/00</u>	a/11	a/11	→a/11	<u>a/00</u>
	frame 1	b/10	b/10	b/10	b/10	b/11	<u>b/00</u>	<u>b/10</u>	<u>b/10</u>	<u>b/10</u>	<u>b/10</u>	d/10
	frame 2	c/10	c/10	c/10	c/10	c/10	e/10	e/10	e/10	e/10	e/10	→e/00
	frame 3	d/10	d/10	d/10	d/10	d/10	→d/00	→d/00	→d/00	→d/00	c/10	c/00
Page fault P2nd-chance Q2nd-chance							F e c				F c d	F d b

- Superscript W : page update
- Underline : the page is on the cleaning list
- P2nd-chance : pages loaded into memory
- Q2nd-chance : pages displaced from memory







# Counting Algorithms

---

- **Keep a counter of the number of references that have been made to each page**
  - Not common
- **Least Frequently Used (LFU) Algorithm**
  - replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm**
  - based on the argument that the page with the smallest count was probably just brought in and has yet to be used

# Workload Examples



# Applications and Page Replacement

---

- All of these algorithms have OS guessing about future page access
- Some applications have better knowledge – i.e. databases
- Memory intensive applications can cause double buffering
  - OS keeps copy of page in memory as I/O buffer
  - Application keeps page in memory for its own work
- Bypasses buffering, locking, etc
- **O\_DIRECT** mode

# Load control strategies

---

- Control the multiprogramming degree of a system
- Related to allocation strategy
- Should maintain multiprogramming degree at the appropriate range
- **Underloaded**
  - System resource waste, performance degradation
- **Overloaded**
  - System resource contention, performance degradation
  - Thrashing (excessive paging activity)
- **Plateau** ranges are located at different positions for different systems

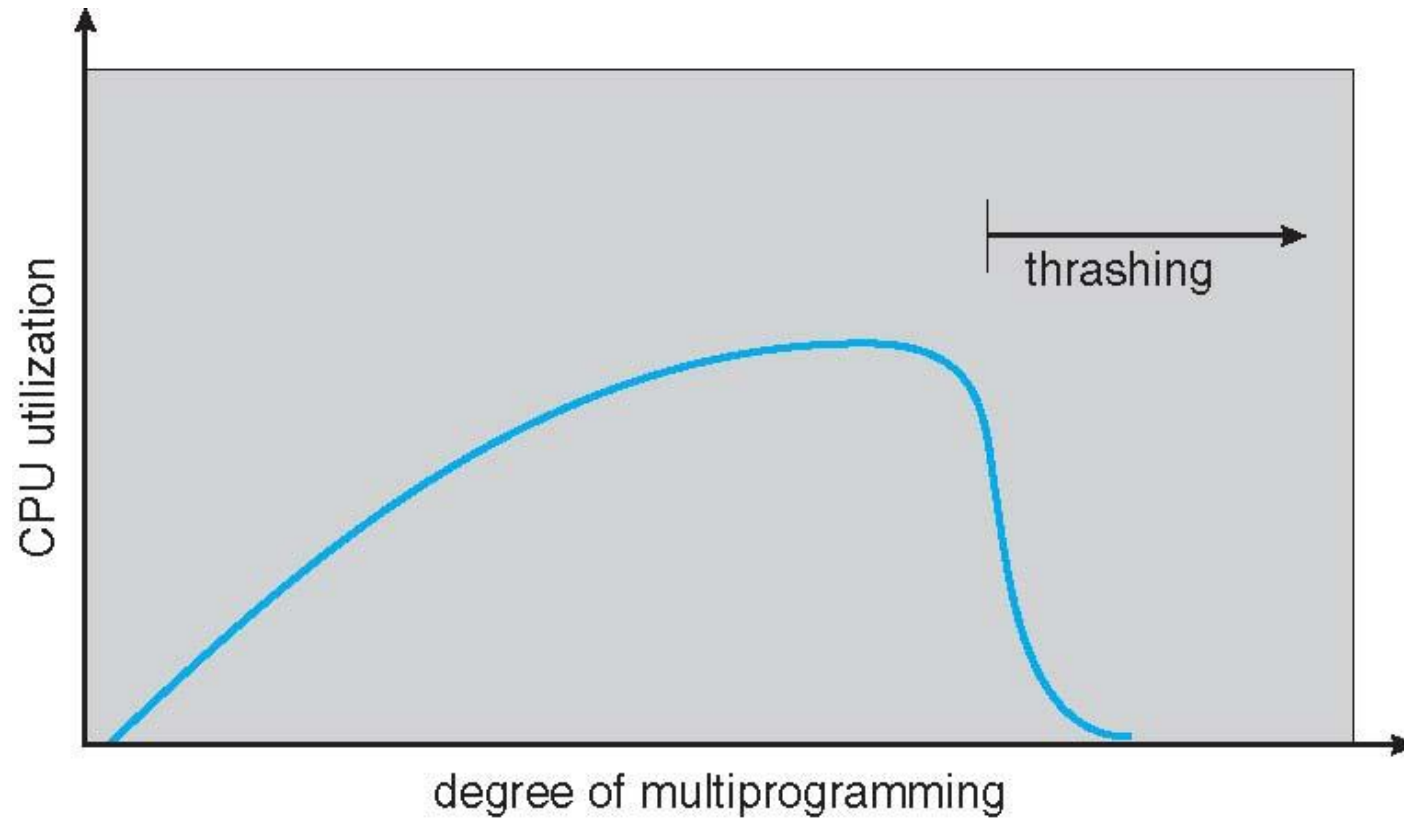
# Thrashing

---

- **If a process does not have “enough” pages, the page-fault rate is very high**
  - Page fault to get page → Replace existing frame → But quickly need replaced frame back
  - This leads to:
    - Low CPU utilization
    - Operating system thinking that it needs to increase the degree of multiprogramming
    - Another process added to the system
- **Thrashing**
  - a process is busy swapping pages in and out

# Thrashing

---



# Demand Paging and Thrashing

---

- **Why does demand paging work?**
  - **Locality model**
    - **Spatial locality and Temporal locality**
  - Process migrates from one locality to another
  - Localities may overlap
- **Why does thrashing occur?**
  - $\Sigma$  size of locality > total memory size
  - Limit effects by using local or priority page replacement

# Be Lazy: Demand Zeroing

---

- To add a page to your address space (heap)
  - OS adds a page to your heap by finding a page in physical memory,
  - zeroing it (required for security)
  - and then mapping it into your address space (i.e., setting up the page table).
  - High cost particularly if the page does not get used
- Demand zeroing
  - OS puts an entry in the page table that marks the page inaccessible
  - If the process then reads or writes the page, a trap takes place.
  - When handling the trap, the OS notices that this is actually a demand-zero page.
  - at this point, the OS then does the needed work of finding a physical page, zeroing it, and mapping it into the process's address space.
  - If the process never accesses the page, all of this work is avoided



# Be Lazy: Copy-on-Write

---

- **Copy-on-Write**

- When OS needs to copy a page from one address space to another, instead of copying it, it can map it into the target address space and mark it read-only in both address spaces. → fast copy
- For read operation, no further action is taken
- For write operation, a trap occurs
  - OS (lazily) allocate a new page, fill it with the data, and map this new page into the address space of the faulting process.

- **fork()** system call


- Creates a child process as a duplicate of its parent
- Should create a copy of the parent's address space for the child, duplicating the pages of the parent

- With **copy-on-write** scheme

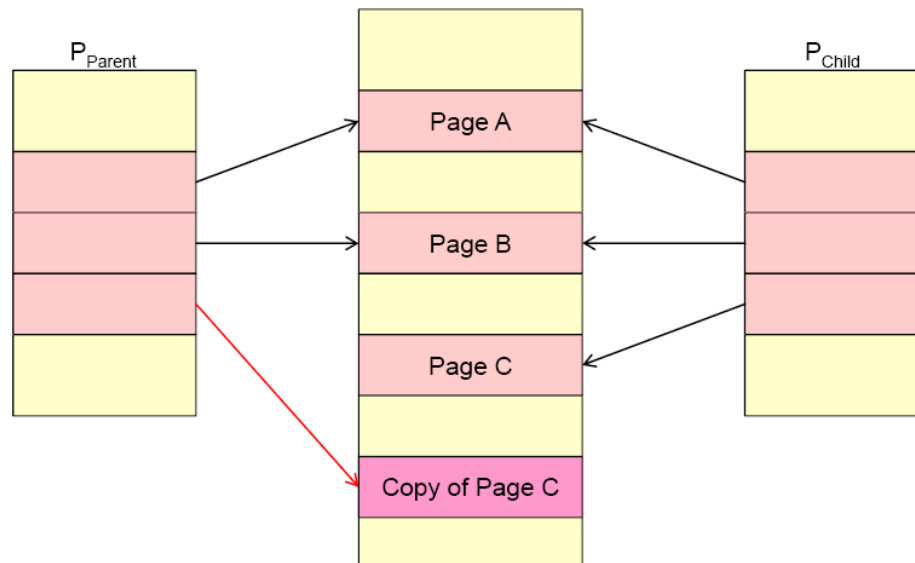
- Allows the parent and the child processes initially share the pages (marked as copy-on-write)
- When any process writes to a shared page, a copy of the shared page is created

# Be Lazy: Copy-on-Write

After **fork()** system call



After Parent modifies Page-C



# Be Lazy: Copy-on-Write

---

- **Notes on copy-on-write**

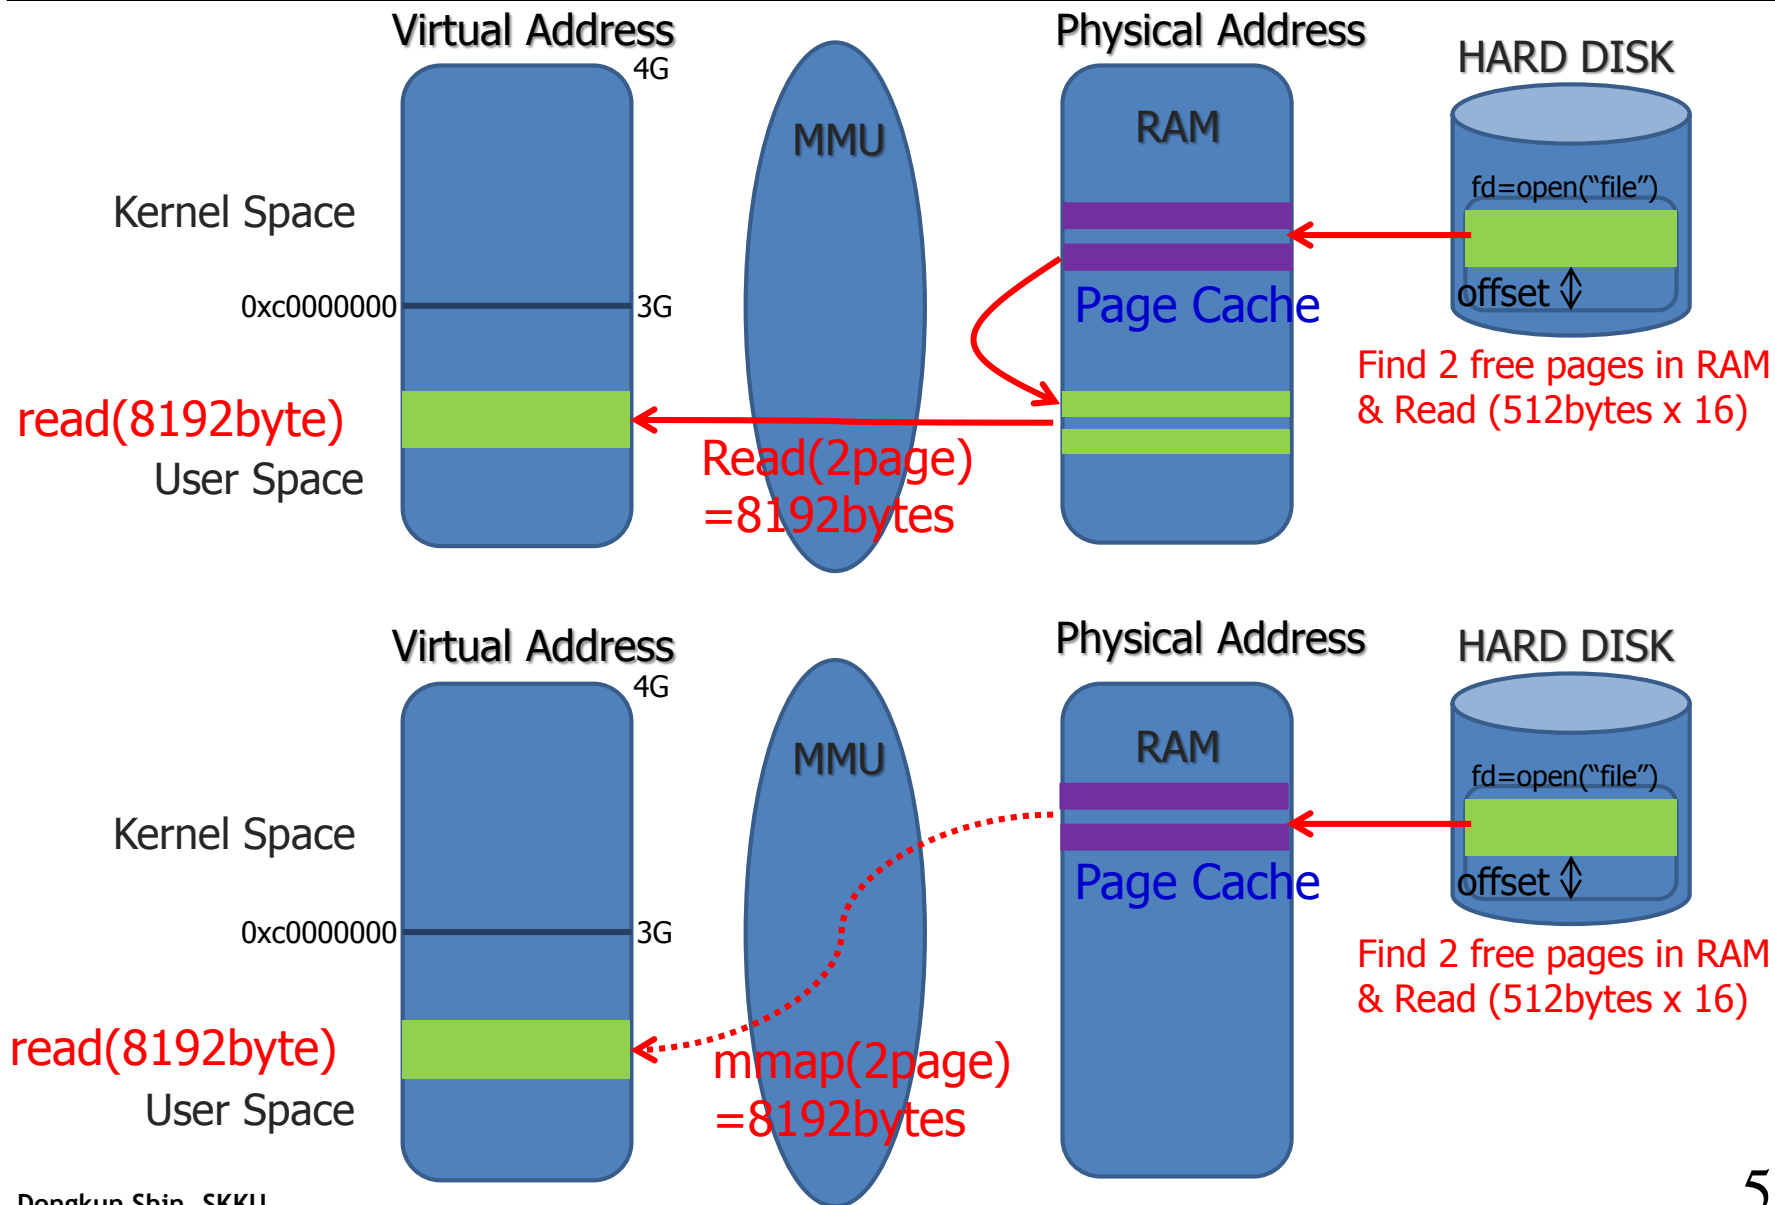
- Only the copy-on-write pages that are modified by either process are copied
  - All unmodified pages are shared by the parent and child processes
- Only the pages that can be modified can be marked as copy-on-write
- Operating systems that use copy-on-write scheme
  - Linux, Solaris, Windows XP

# Memory-Mapped Files

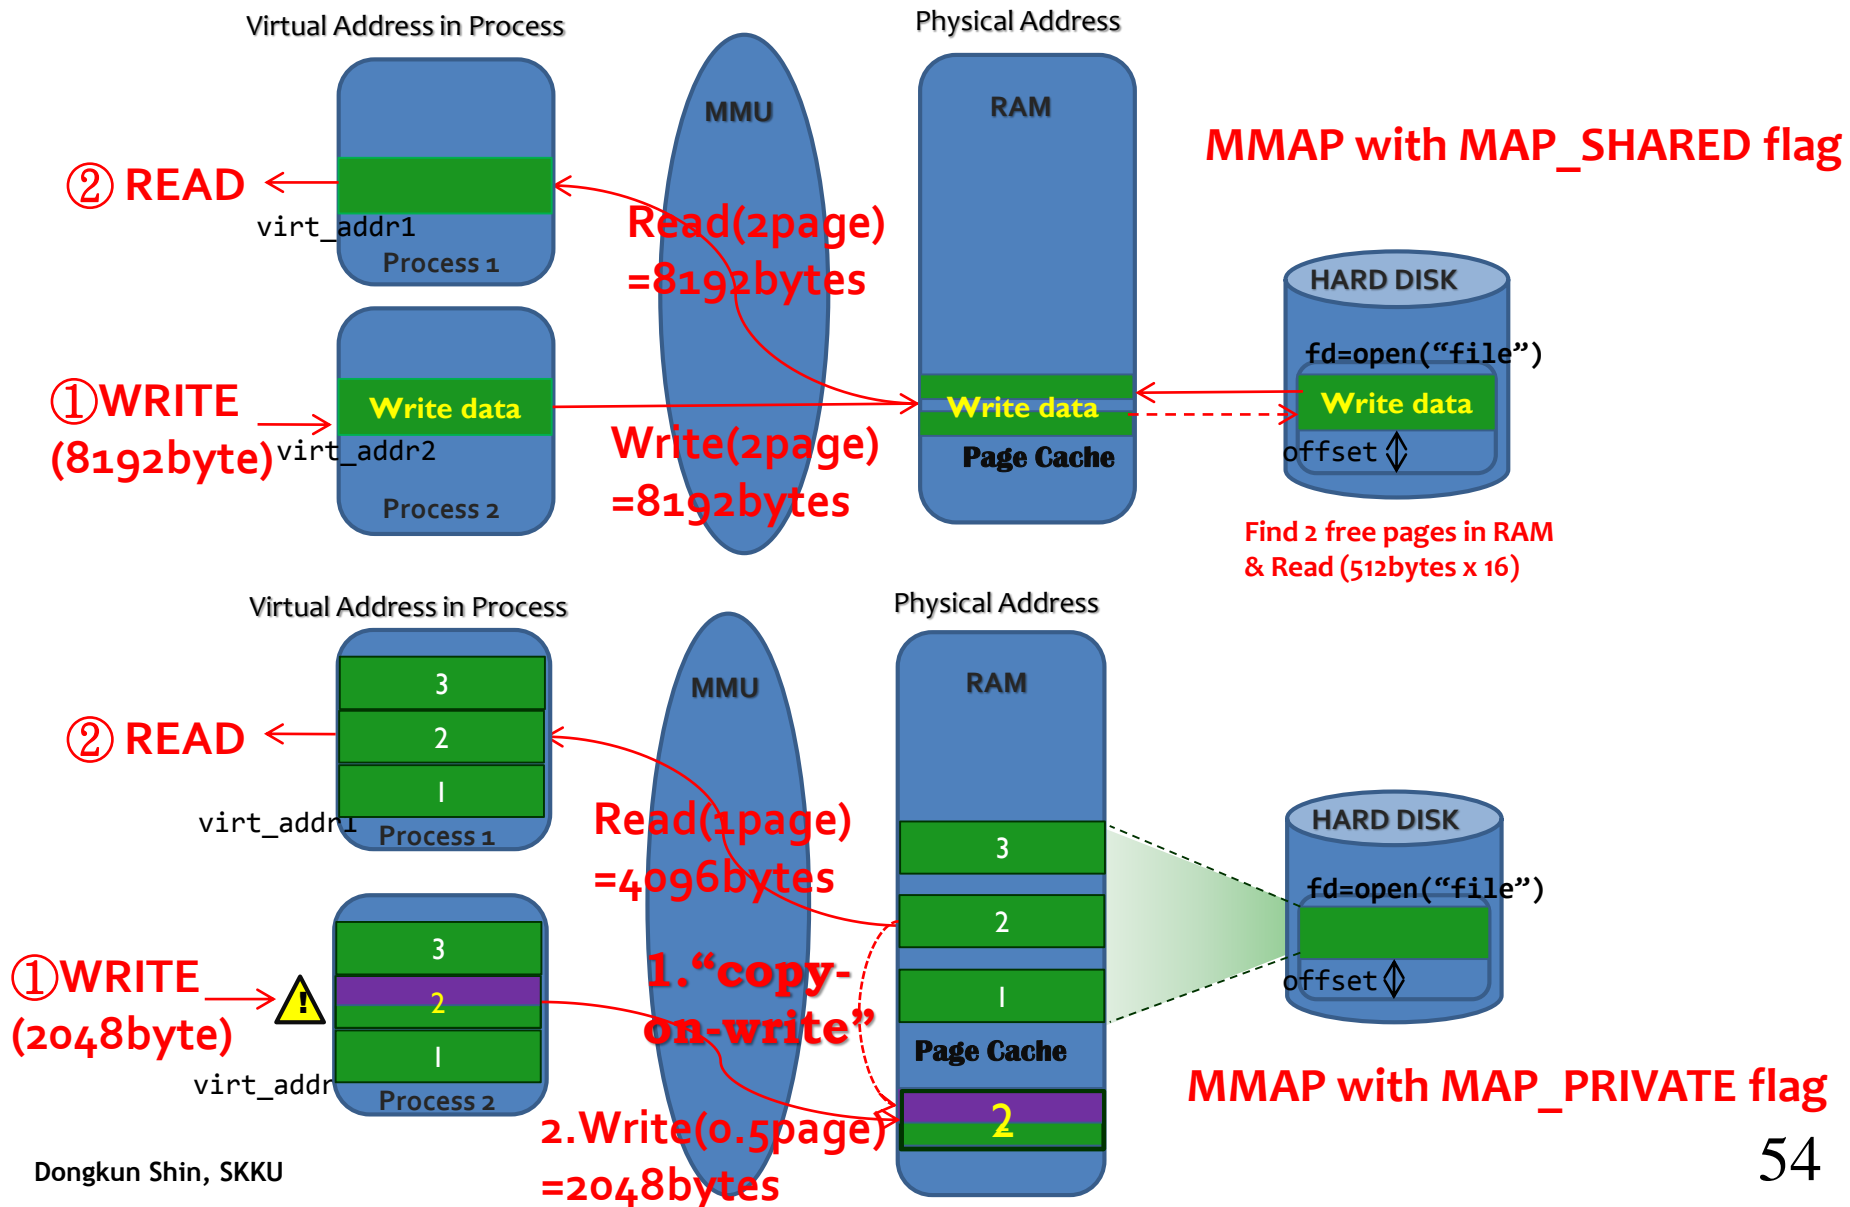
---

- **Memory-mapped file I/O allows file I/O to be treated as routine memory access by `mapping` a disk block to a page in memory**
- **A file is initially read using demand paging**
  - A page-sized portion of the file is read from the file system into a physical page
  - Subsequent reads/writes to/from the file are treated as ordinary memory accesses
- **Simplifies and speeds file access by driving file I/O through memory rather than `read()` and `write()` system calls**
- **Also allows several processes to map the same file allowing the pages in memory to be shared**
- **But when does written data make it to disk?**
  - Periodically and / or at file `close()` time
  - For example, when the pager scans for dirty pages

# MMAP



# MMAP on Shared File



# Other Considerations -- Prepaging

---

- **Prepaging (Prefetch)**

- To reduce the large number of page faults that occurs at process startup
- Prepage all or some of the pages a process will need, before they are referenced
- But if prepaged pages are unused, I/O and memory was wasted
- Assume  $s$  pages are prepaged and  $a$  of the pages is used
  - Is cost of  $s * a$  save pages faults  $>$  or  $<$  than the cost of prepaging  $s * (1 - a)$  unnecessary pages?
  - $a$  near zero  $\Rightarrow$  prepaging loses

# Other Issues – Page Size

---

- **Sometimes OS designers have a choice**
  - Especially if running on custom-built CPU
- **Page size selection must take into consideration:**

- Fragmentation
- Page table size
- Resolution
- I/O overhead
- Number of page faults
- Locality
- TLB size and effectiveness

## Smaller page size

- + Smaller internal fragmentation
- + Match program locality more accurately
  - o Reduction in total I/O
  - o Better utilization of memory  
(Less total allocation of memory)
- Larger page table size
- Increase in I/O time
- Increase in the number of page faults

- **Always power of 2, usually in the range  $2^7$  (128 bytes) to  $2^{22}$  (4MB)**
- **On average, growing over time**



# Other Issues – TLB Reach (Coverage)

---

- **TLB Reach**
  - The amount of memory accessible from the TLB
  - $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- **Ideally, the working set of each process is stored in the TLB**
  - Otherwise there is a high degree of page faults
- **Increase the number of entries in TLB**
  - Expensive
- **Increase the Page Size**
  - lead to an increase in fragmentation as not all applications require a large page size
- **Provide Multiple Page Sizes**
  - Allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation
  - Requires OS support
    - One of the fields in TLB entry must indicate the size of the page (frame) corresponding to the TLB entry

# Other Issues – Program Structure

---

- **Program restructuring**
  - System performance can be improved if the user (or compiler/linker/loader) has an awareness of the paged nature of memory or underlying demand paging system
  - Program restructuring by user or compiler/linker/loader

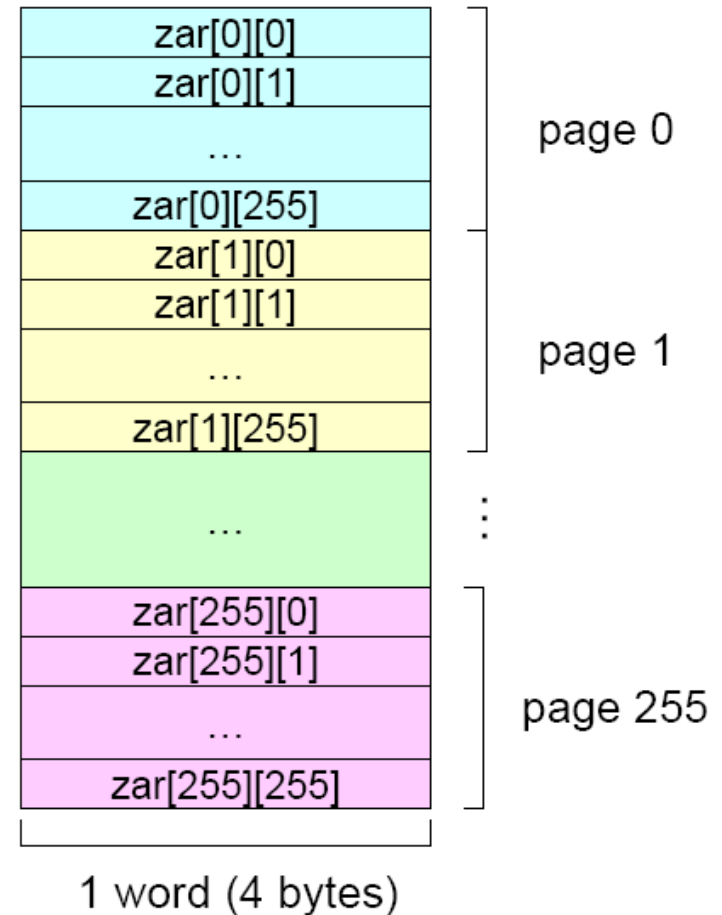
# Other Issues – Program Structure

- **Example**

- Assumptions
  - Paging system of 1KB page size
  - sizeof(int): 4 Bytes

```
// Program-1
int main()
{
    int zar[256][256];
    int i, j;

    for(j = 0; j < 256; j++)
        for(i = 0; i < 256; i++)
            zar[i][j] = 0;
    return 0;
}
```



For each execution of loop instance, pages 0~255 are referenced sequentially and this process is repeated 255 times

For execution with no page faults, The process should have 256 page frames


- Not practical

# Other Issues – Program Structure

---

```
// Program-1
int main()
{
    int zar[256][256];
    int i, j;

    for(j = 0; j < 256; j++)
        for(i = 0; i < 256; i++)
            zar[i][j] = 0;
    return 0;
}
```



```
// Program-2
int main()
{
    int zar[256][256];
    int i, j;

    for(i = 0; i < 256; i++)
        for(j = 0; j < 256; j++)
            zar[i][j] = 0;
    return 0;
}
```

# Other Issues – Program Structure


---

- Notes
  - User
    - Careful selection of data structures and programming structures can increase locality (hence lower page-fault rate)
      - Stack: has good locality
      - Hash table: produces bad locality
      - Pointer: diminishes locality
    - OOPs tend to have a poor locality of reference
  - Linker/Loader
    - Avoid placing routines across page boundaries
    - Let routines that call each other many times be packed into the same page

# Other Issues – I/O interlock

---

- **I/O Interlock** – Pages must sometimes be locked into memory
  - Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm (lock bit)
  - Another solution: never to execute I/O to user memory
- **Pinning of pages to lock into memory**



# Homework

---

- Homework in Chap 21 (**vmstat**)
- Homework in Chap 22 (paging-policy.py)