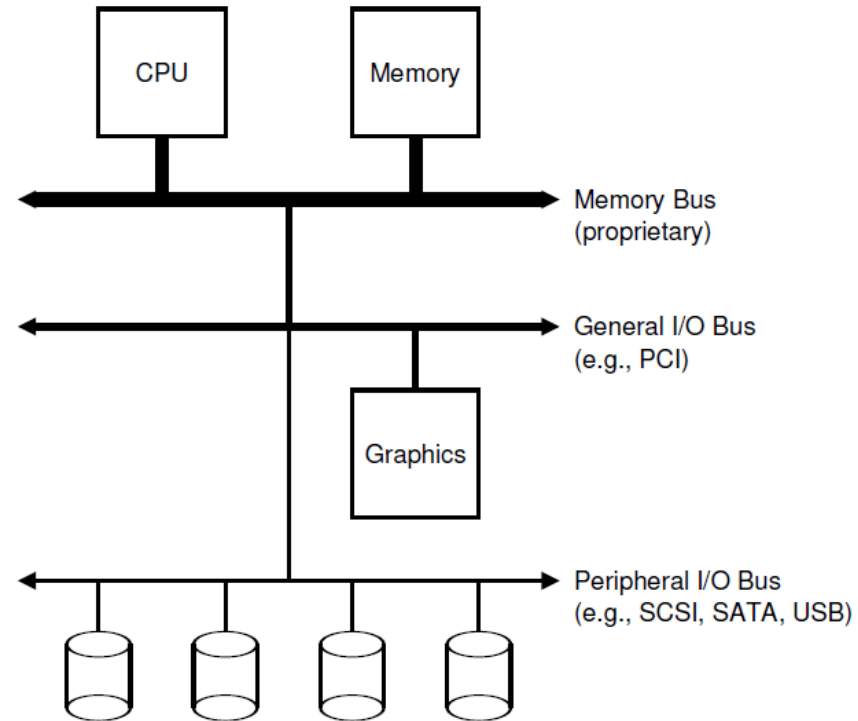

I/O Devices & SSD

System Architecture

- **Hierarchical approach**
- **Memory bus**
 - CPU and memory
 - Fastest
- **I/O bus**
 - e.g., PCI
 - Graphics and higher-performance I/O devices
- **Peripheral bus**
 - SCSI, SATA, or USB
 - Connect many slowest devices



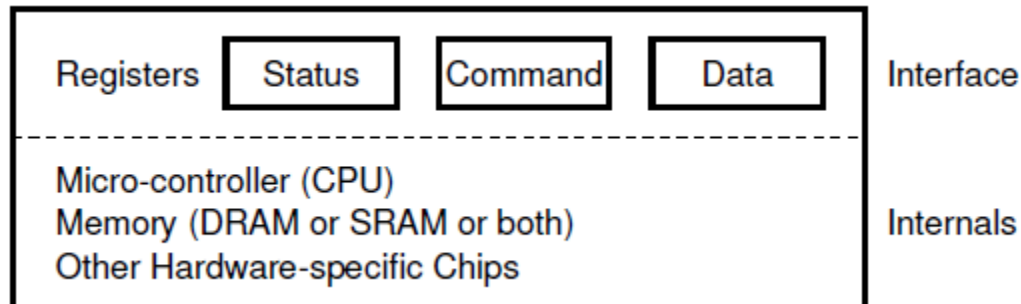
Device Structure

- **Hardware Interface**

- Allows the system software to control its operation
- status register: can be read to see the current status of the device
- command register: can be written to tell the device to perform a certain task
- data register: transfer data to/from the device

- **Internal Structure**

- Implementation specific
- Simple devices
 - have one or a few hardware chips to implement their functionality;
- Complex devices
 - include a simple CPU running **firmware**, some general purpose memory, and other device-specific chips



Protocol

- By reading and writing internal registers of a device, the operating system can control device behavior.

```
While (STATUS == BUSY)
    ; // wait until device is not busy -> polling
Write data to DATA register
Write command to COMMAND register
    (Doing so starts the device and executes the command)
While (STATUS == BUSY)
    ; // wait until device is done with your request
```

- Inefficiencies and Inconveniences
 - Polling
 - Repeatedly reading the status register
 - Programmed I/O (PIO)
 - The main CPU is involved with the data movement

Lowering CPU Overhead With Interrupts

- **Interrupt**

- OS can issue a request, put the calling process to sleep, and context switch to another task.
 - When the device is finally finished with the operation, it will raise a hardware interrupt, causing the CPU to jump into the OS at a pre-determined interrupt service routine (ISR)
 - The handler in operating system code that will finish the request
 - e.g., reading data and perhaps an error code from the device
 - Wake the process waiting for the I/O
 - **Allow for overlap of computation and I/O**
- Interrupt is not *always* the best solution
 - If a device is fast, it may be best to poll.
 - Otherwise, use interrupt

Interrupt Optimization

- **Unknown device: hybrid approach**

- polls for a little while and then, if the device is not yet finished, uses interrupts. (two-phased approach)

- **Network systems**


- When a huge stream of incoming packets each generate an interrupt, it is possible for the OS to **livelock**
 - only processing interrupts and never allowing a user-level process to run
- Occasionally use polling to better control what is happening in the system (e.g., slashdot effect)

- **Interrupt Coalescing**


- Multiple interrupts are coalesced into a single interrupt delivery
 - lower the overhead of interrupt processing
- Long waiting
 - Many interrupts can be coalesced
 - Increase the latency of a request

More Efficient Data Movement With DMA

- OS programs the DMA engine
 - Source/destination address in memory or device
 - Data amount
- Other processes can be serviced
- When the DMA is complete, the DMA controller raises an interrupt, and the OS thus knows the transfer is complete.



Programmed IO



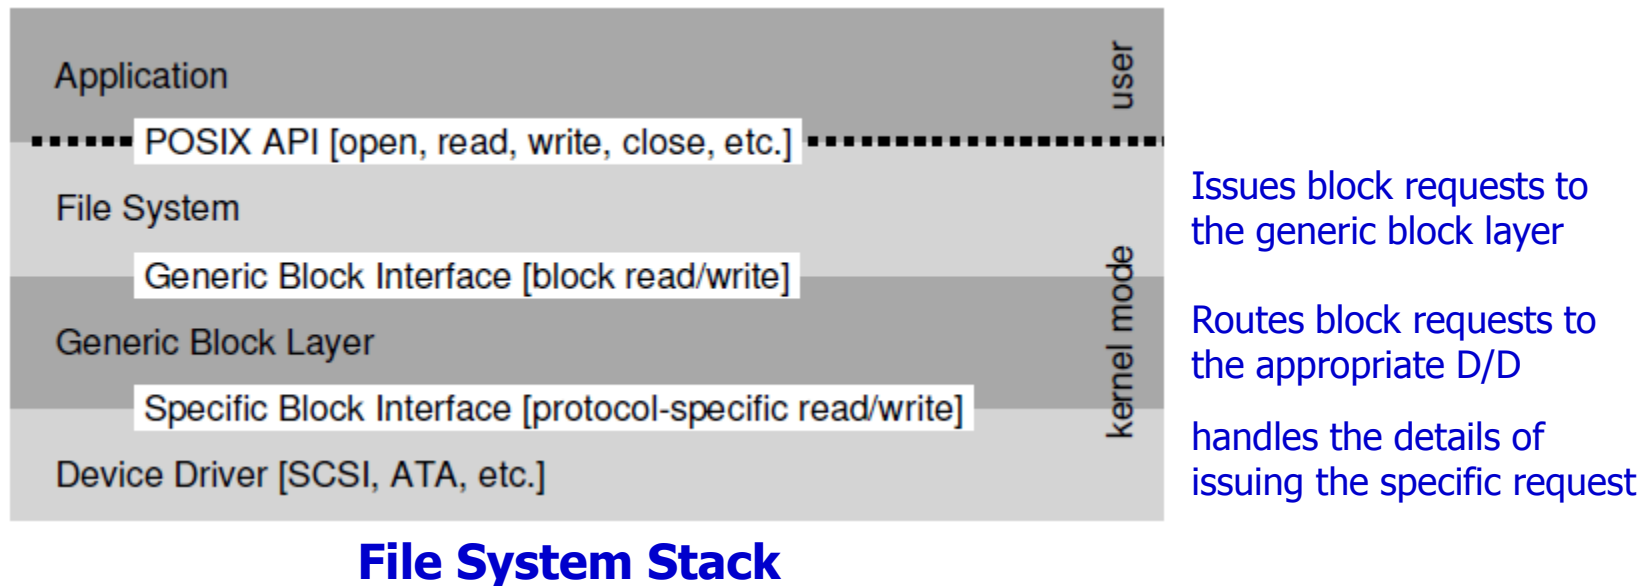
DMA

Methods Of Device Interaction

- **I/O instructions**
 - E.g., in and out on x86
 - To send data to a device, register → **port**.
 - Usually **privileged** instructions
- **Memory-mapped I/O**
 - Device registers are mapped to memory address space
 - load (to read) or store (to write)

Device Driver


- How can we keep most of the OS device-neutral, thus hiding the details of device interactions from major OS subsystems?
- Solution: **Abstraction**
 - The device driver in the OS know in detail how a device works
 - Any specifics of device interaction are encapsulated within D/D



Device Driver

- Drawbacks of encapsulation
 - Even if a device has many special capabilities, it has to present a generic interface to the rest of the kernel, those special capabilities will go unused.
 - e.g., SCSI devices which have very rich error reporting
- Over 70% of OS code is found in device drivers
 - for any given installation, most of that code may not be active
 - have many more bugs and thus are a primary contributor to kernel crashes

Disk System



Disk System

- **Average time to access some target sector**

- $T_{\text{access}} = T_{\text{avg-seek}} + T_{\text{avg-rotation}} + T_{\text{avg-transfer}}$

- Seek time ($T_{\text{avg-seek}}$)

- Time to position heads over cylinder containing target sector

- Typical $T_{\text{avg-seek}}$ is 3~9ms

- Rotational latency ($T_{\text{avg-rotation}}$)

- Time waiting for first bit of target sector to pass under head

- $T_{\text{avg-rotation}} = 1/2 \times 1/\text{RPMs} \times 60\text{s}/1\text{min}$

- Typical $T_{\text{avg-rotation}} = 1\sim 4\text{ms}$

- Transfer time ($T_{\text{avg-transfer}}$)


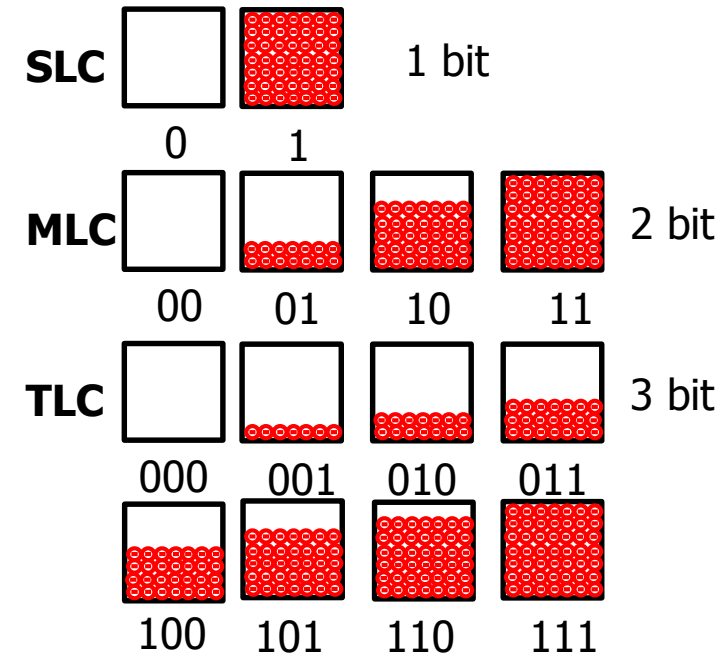
- Time to read the bits in the target sector

- $T_{\text{avg-transfer}} = 1/\text{RPM} \times 1/(\text{avg. \# sectors/track}) \times 60\text{s}/1\text{min}$

Flash-based SSD

- **NAND flash memory**

- Page: read/write unit, 4KB~16KB
- Block: erase unit, **erase-before-write**, 128 KB or 256 KB
- Limited lifetime: wear out
- SLC, MLC, TLC



Flash Operations

- **Read (a page)**
 - Random access
- **Erase (a block)**
 - set each bit to the value 1
- **Program (a page)**
 - change some of the 1's within a page to 0's

Device	Read (μ s)	Program (μ s)	Erase (μ s)
SLC	25	200-300	1500-2000
MLC	50	600-900	~3000
TLC	~75	~900-1350	~4500

Page 0	Page 1	Page 2	Page 3
00011000	11001110	00000001	00111111
VALID	VALID	VALID	VALID

Page 0	Page 1	Page 2	Page 3
11111111	11111111	11111111	11111111
ERASED	ERASED	ERASED	ERASED

Page 0	Page 1	Page 2	Page 3
00000011	11111111	11111111	11111111
VALID	ERASED	ERASED	ERASED

Erase Block

Program Page 0

Previous contents of Page 1, 2, 3 are all gone!
They must be moved before the block erase operation.


Flash Performance And Reliability

- **Wear out**

- Erased/program slowly accrues a little bit of extra charge.
- As that extra charge builds up, the block becomes unusable.
- MLC block: 10,000 P/E (Program/Erase) cycle lifetime;
- SLC block: 100,000 P/E cycles.


- **Disturbance**

- When accessing a particular page within a flash, it is possible that some bits get flipped in neighboring pages
- read disturbs or program disturbs





Program
F-N Tunneling

Off cell
(Solid-0)



Erase
F-N Tunneling

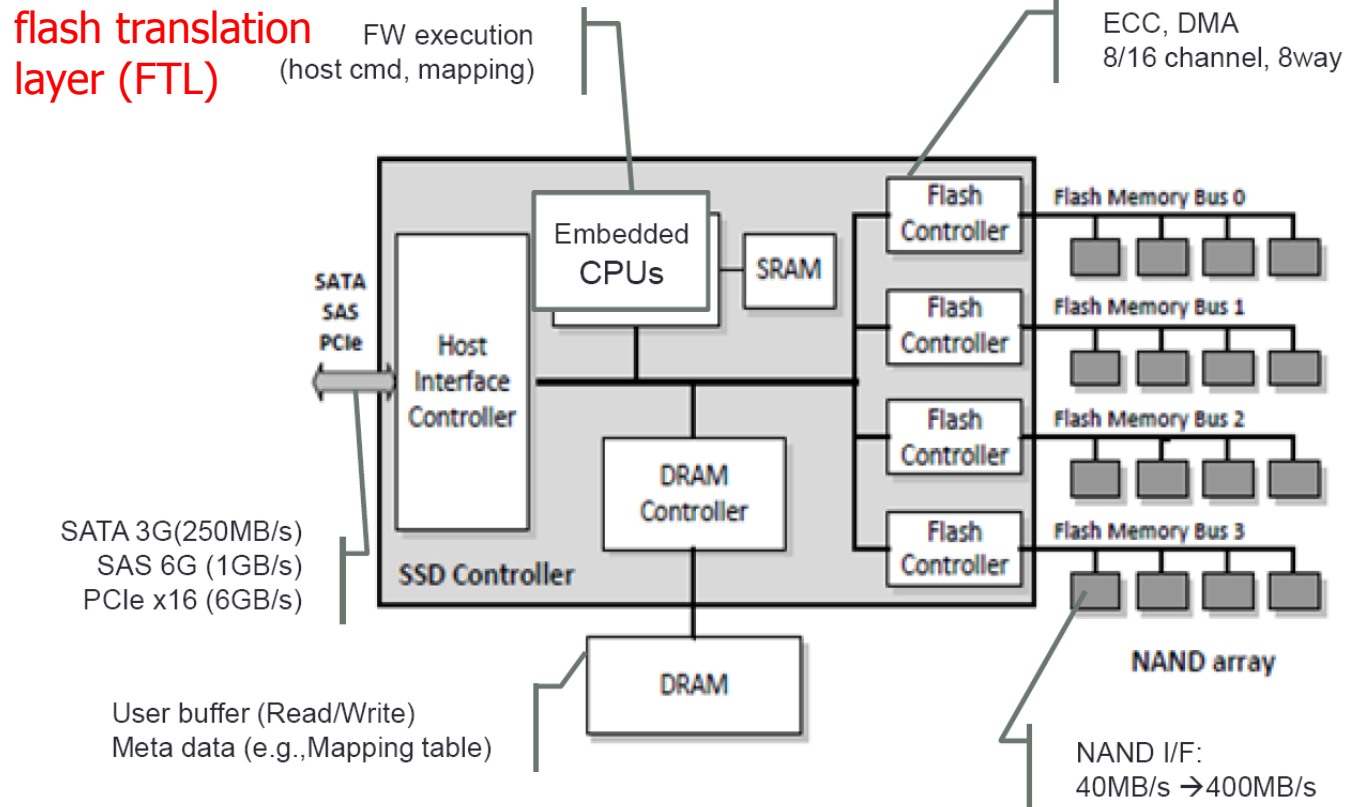
On cell
(Solid-1)




cell-to-cell interference

From Raw Flash to Flash-Based SSDs

- Flash-based SSD provides standard **block interface** atop the raw flash chips inside it.
 - Sector (512 KB) read/write



Samsung Galaxy Note8



Solid-State Disk




SATA SSD





NVMe SSD



M.2 NVMe



MacBook Air



- FTL takes read and write requests on **logical blocks** and turns them into low-level read, erase, and program commands on the underlying **physical blocks** and **physical pages**
- Performance issues
 - Utilizing multiple flash chips in parallel
 - Reduce write amplification
 - total write issued to flash chips/total write issued by host
- Reliability issues
 - Wear leveling
 - spread writes across the blocks of the flash as evenly as possible
 - ensure that all of the blocks of the device wear out at roughly the same time
 - Program disturbance
 - sequential-programming
 - program pages within an erased block in order
 - low page → high page


FTL: Direct Mapped

- **A write to logical page N**
 - Read in the entire block that page N is contained within
 - Erase the block
 - Program the old pages as well as the new one.
 - Read-modify-write approach
- **Poor performance & high write amplification**
- **The physical blocks containing popular data will quickly wear out.**


A Log-Structured FTL

- Upon a write to logical block N
 - The device appends the write to the next free spot in the currently-being-written-to block
 - no overwrite, out-of-place update
 - To allow for subsequent reads of block N, the device keeps a **mapping table**, which stores the physical address of each logical block.
- Problems: GC & mapping table


① Initial page address mapping for page (B)



② Update data (write new data) in an empty page



③ Page address re-mapping for page (B)




Where is the mapping table?
What happens if the device loses power?


Garbage Collection

- dead-block reclamation to provide free space for new writes
 - find a block that contains one or more garbage pages
 - read in the live (non-garbage) pages from that block, write out those live pages to the log
 - (finally) reclaim the entire block for use in writing
- Must know
 - Whether each page is live or dead → mapping table or bitmap
 - The number of live pages in each block for victim selection
- Overprovision can delay GC, and push to the background


① Number of empty blocks becomes short




② Collect (copy) valid pages into an empty block



③ Page address re-mapping

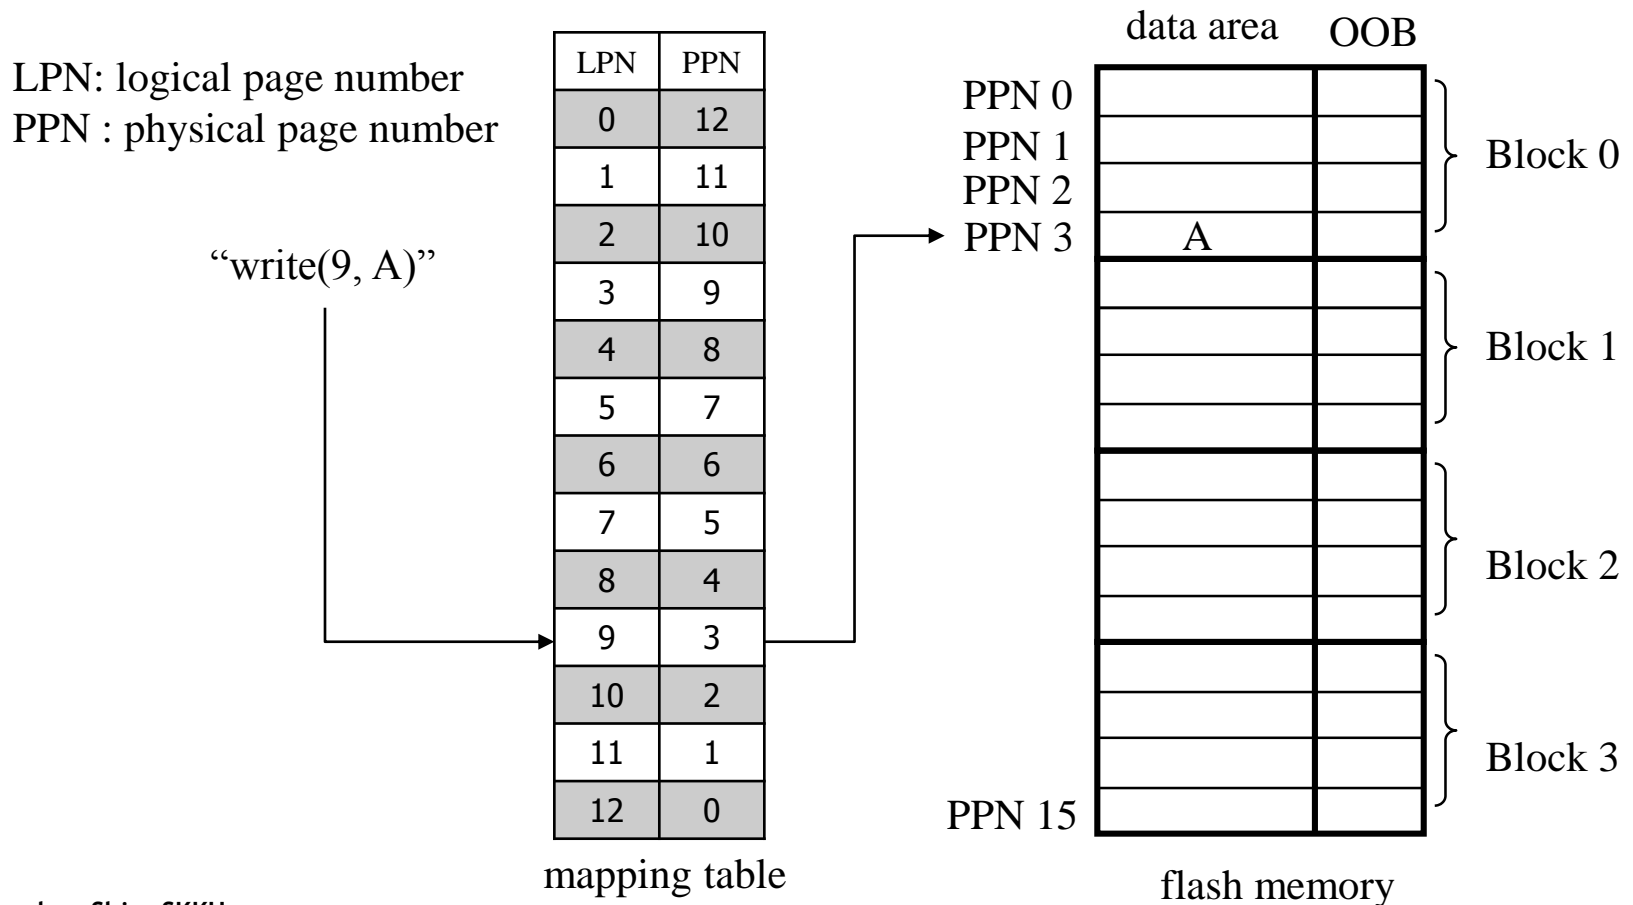


④ Erase invalid blocks (can be re-used now)



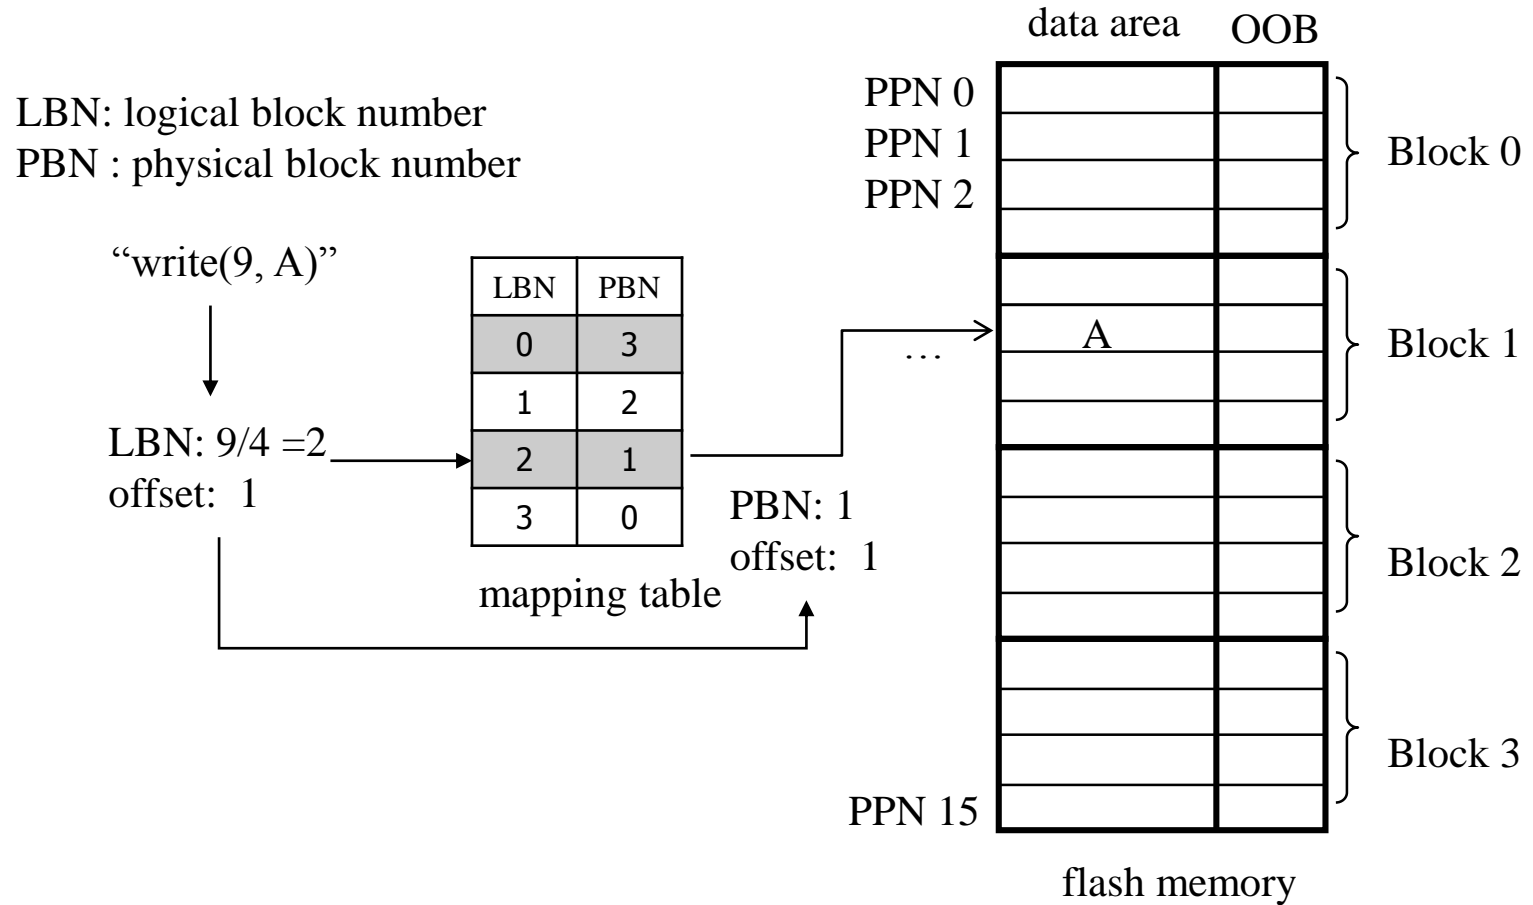
Mapping Table Size

- Page-level mapping
 - With a large 1-TB SSD, a single 4-byte entry per 4-KB page results in 1 GB of memory needed the device




Mapping Table Size

- Block-level mapping




Update at Block Mapping


① Initial block address mapping




② Update data (write new data) in an empty block



③ Block address re-mapping



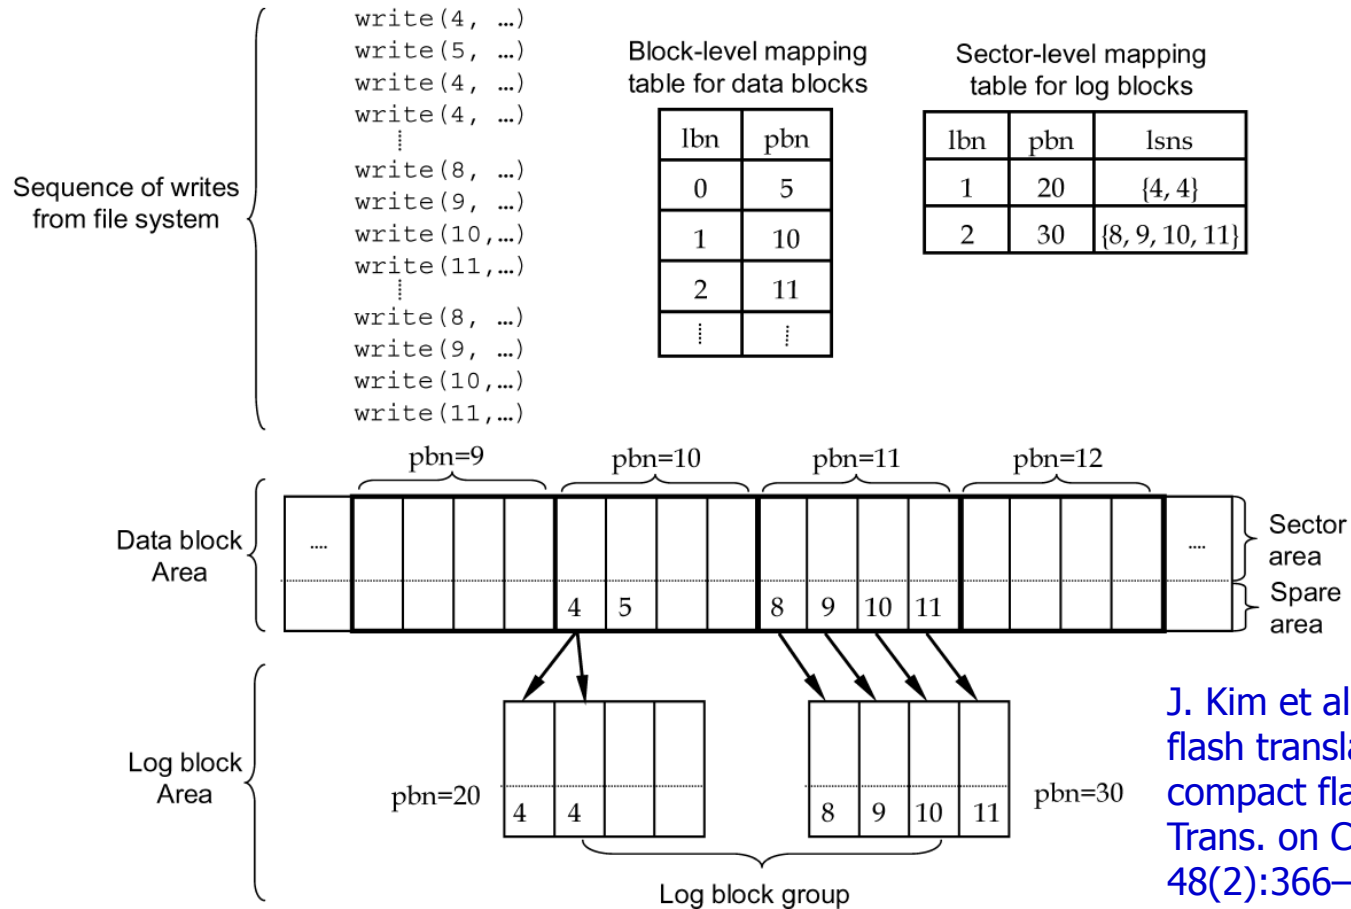
④ Erase the original block (can be re-used now)



Hybrid Mapping

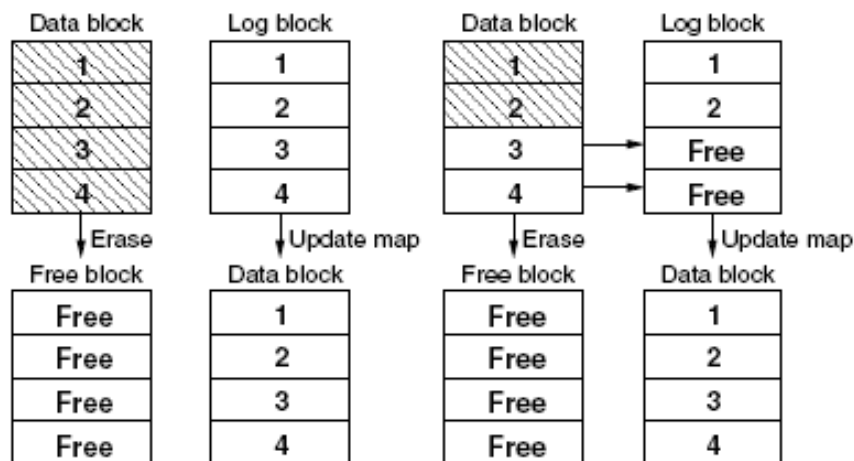
- To enable flexible writing but also reduce mapping costs
- FTL keeps a few **log blocks** and directs all writes to them
 - Page-level mapping (log table)
 - Small number of log blocks → small-sized log table
- Normal data blocks
 - Block-level mapping (data table) → small-sized data table
- For read request
 - first consult the log table; if not found, consult the data table
- To keep the number of log blocks small, log blocks must be merged with data blocks if no free space is available in log blocks.

Hybrid Mapping (BAST)




Hybrid Mapping – Merge Operation

Number Valid page
 Number Invalid page
 Free Free page

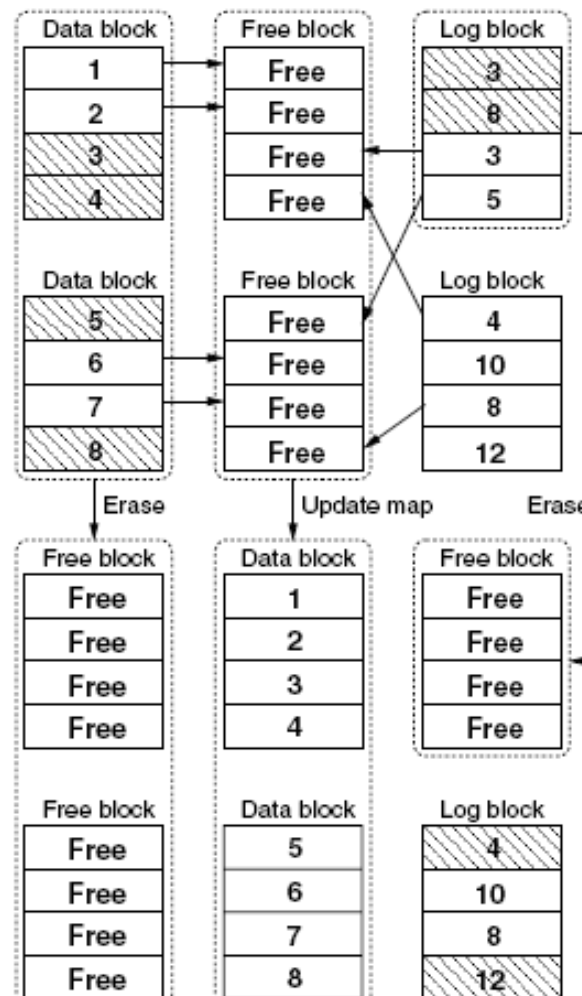


(a) switch merge

(b) partial merge

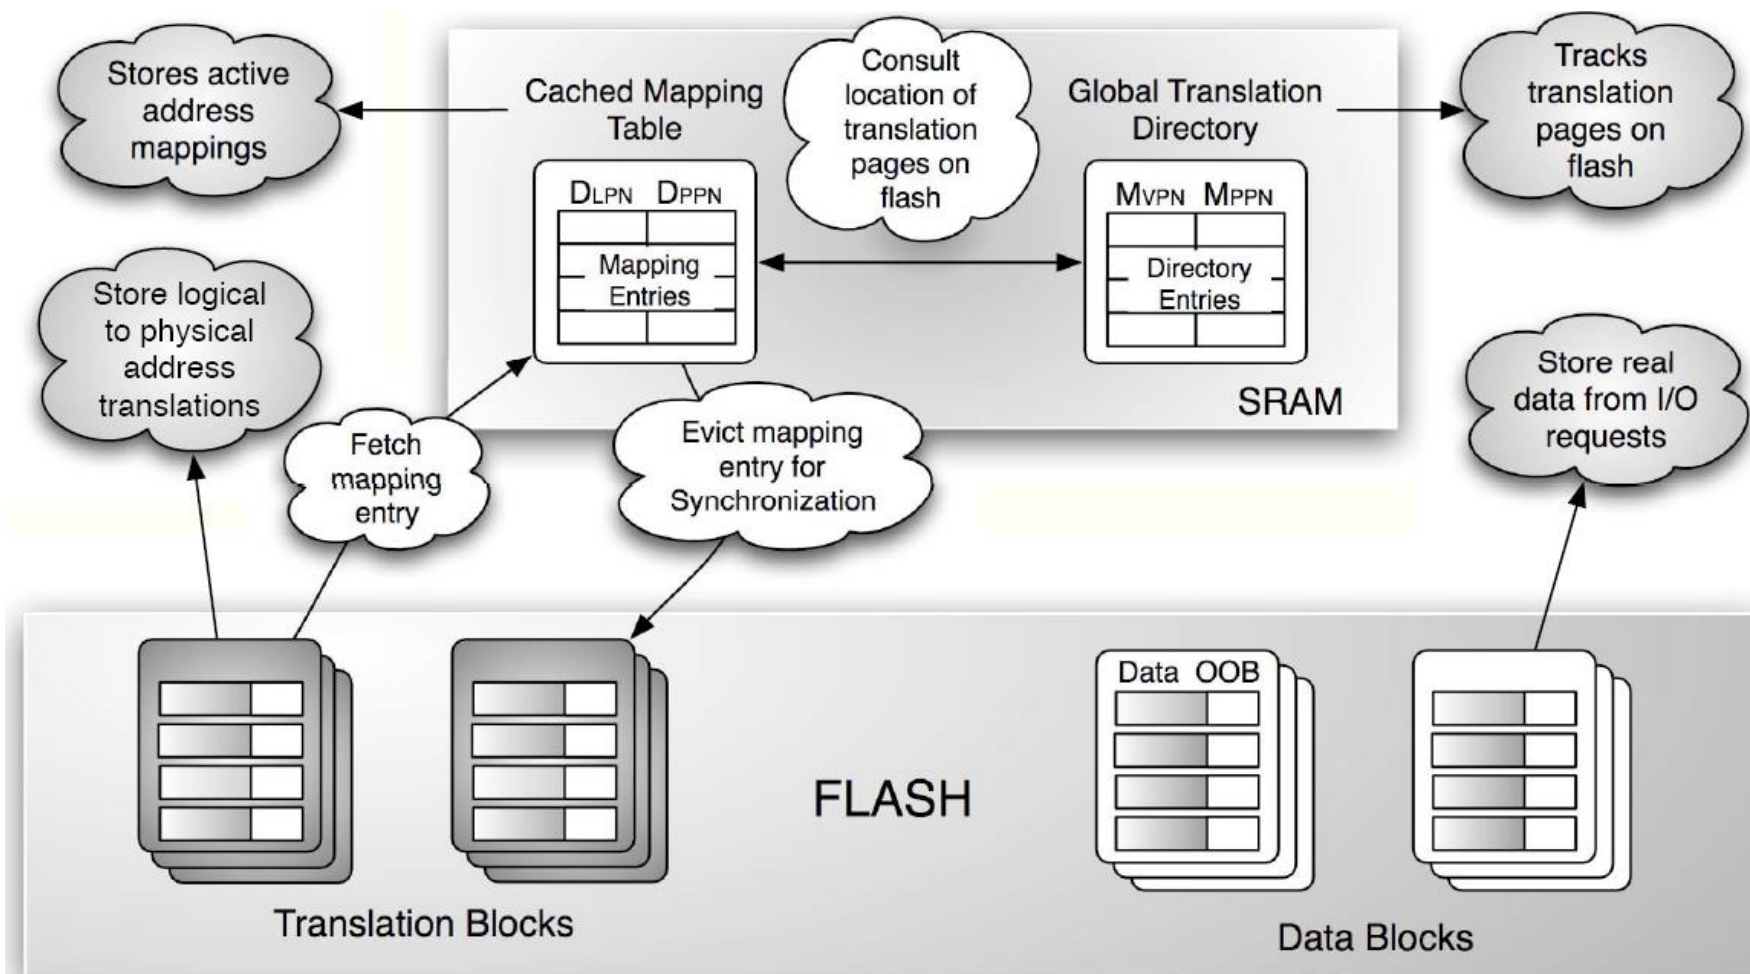


(c) full merge in BAST



(d) full merge in FAST

DFTL



SSD Performance

- Random I/O
 - Large difference between SSD and HDD
- Sequential I/O
 - Much less of a difference between SSD and HDD
 - HDD still a good choice
- SSD random read performance is not as good as SSD random write performance
 - Write buffer
- Performance difference between sequential and random in SSD
 - many of the file system techniques for HDDs are still applicable to SSDs

Device	Random		Sequential	
	Reads (MB/s)	Writes (MB/s)	Reads (MB/s)	Writes (MB/s)
Samsung 840 Pro SSD	103	287	421	384
Seagate 600 SSD	84	252	424	374
Intel SSD 335 SSD	39	222	344	354
Seagate Savvio 15K.3 HDD	2	2	223	223